

# Solving Combinatorial Optimization Problems with Graph Networks

Defeng Liu

Supervised by Andrea Lodi

2020.02.24



**POLYTECHNIQUE  
MONTRÉAL**



**DATA SCIENCE  
FOR REAL-TIME  
DECISION-MAKING**

# Outline

- ▶ Introduction
- ▶ ML for combinatorial optimization
- ▶ Applications with GNNs

# Outline

- ▶ **Introduction**
- ▶ ML for combinatorial optimization
- ▶ Applications with GNNs

# Combinatorial Optimization

A large range of real world decision problems are **Combinatorial Optimization(CO)** problems:

- Examples: supply chain, packing and scheduling, etc...
- Techniques: branch and bound, cutting planes, heuristics, pre-processing ...



The Traveling Salesman Problem  
[examples.gurobi.com](http://examples.gurobi.com)



Loading a 20-foot shipping container  
[www.movehub.com](http://www.movehub.com)

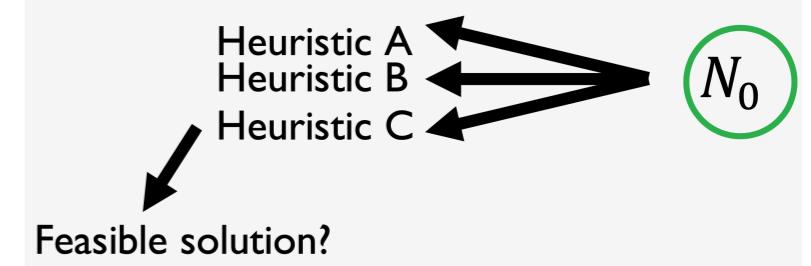
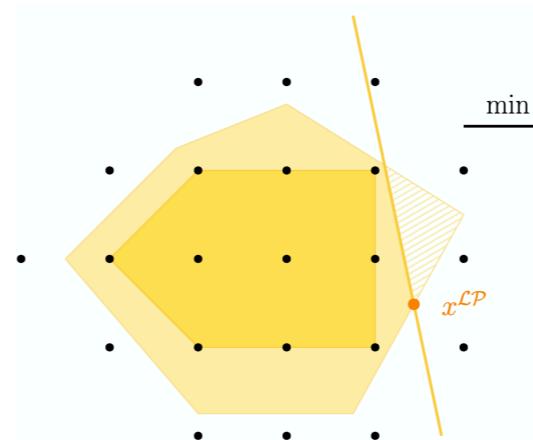
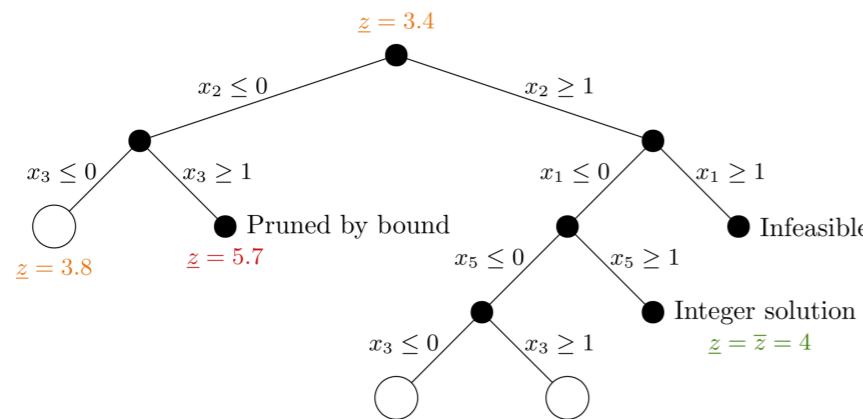
# Combinatorial Optimization

A large range of real world decision problems are **Combinatorial Optimization(CO)** problems:

- Examples: supply chain, packing and scheduling, etc...
- Techniques: branch and bound, cutting planes, heuristics, pre-processing ...

Many of them are NP-complete, practically we develop different methods:

- Exact methods, eg. branch and bound, cutting planes
- Approximation methods
- Heuristics, meta-heuristics

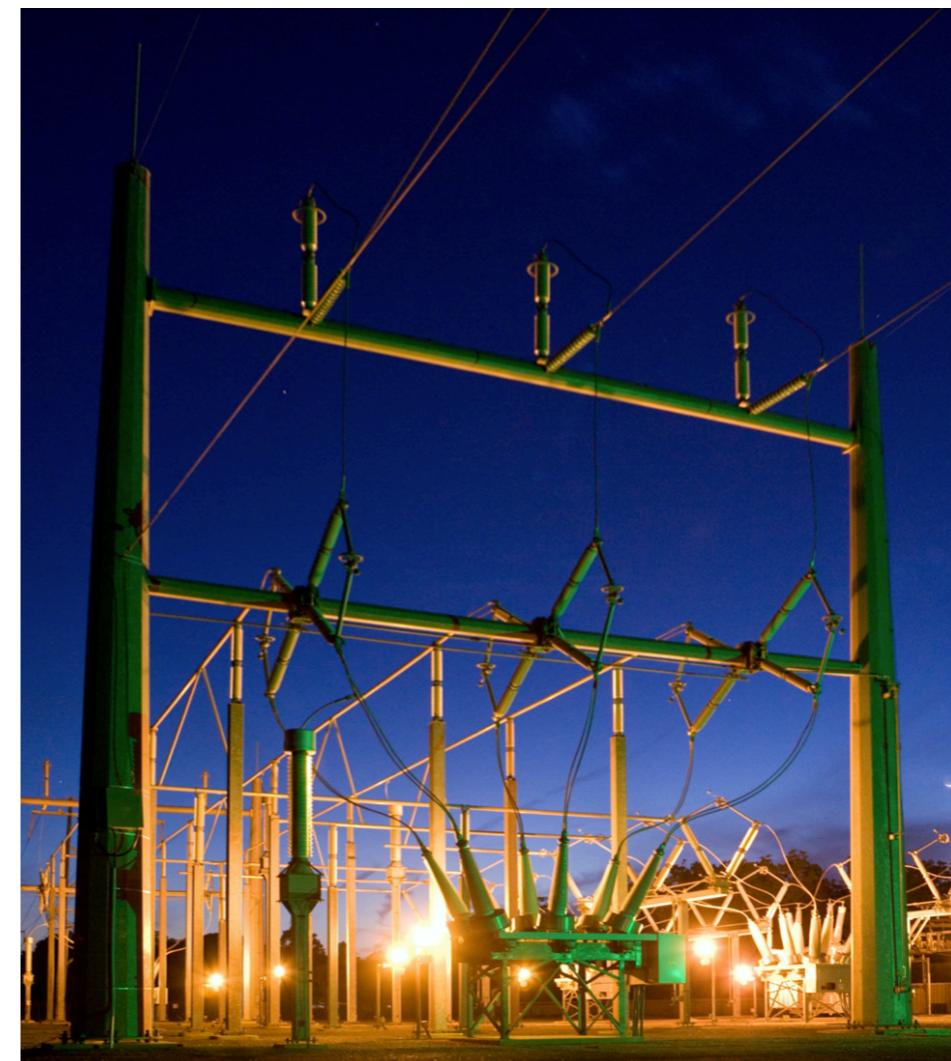


# Motivation

- Although we have sophisticated solvers(*CPLEX*, *Gurobi*, etc), solving CO problems is still a challenging computation process:
  - For CO problems:
    - Challenge1: restricted computation/time budget
      - ▶ Desirable to compute fast solutions
  - For CO algorithms:
    - Challenge2: complex and heavy auxiliary tasks, ex. branching strategies, cuts selection
      - ▶ Cheaper approximators / new decision policies

# Motivation

- In many real world applications, instances of the same problem are solved repeatedly.
  - Application in power systems [Xavier et al. 2019]
    - Schedule 3.8 trillion kWh production ( \$400 billion market annually in US )
    - Solved multiple times daily
    - 10x speed up combining ML and MIP



# Hypothesis - What can be learned from data?

- In state space of CO problems, the instances in neighborhoods
  - have inherent structural commonalities,
  - show similar solving procedures.
  - their solutions are close to each other in solution space
- Valuable patterns/characteristics can be learned from data.
  - approximations of probability distributions of random elements
  - solutions at different levels of detail
  - decision policies in CO algorithms

# Requirement

- We want to keep the guarantees provided by exact OR algorithms (**feasibility, optimality**)

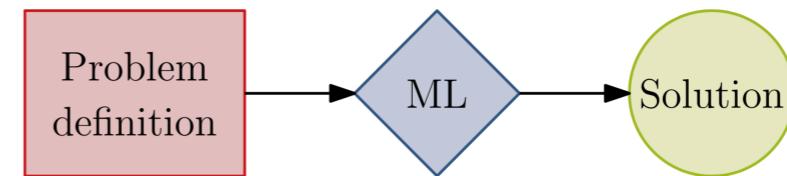
# Outline

- ▶ Introduction
- ▶ **ML for combinatorial optimization**
- ▶ Applications with GNNs

# Methodology - Paradigms

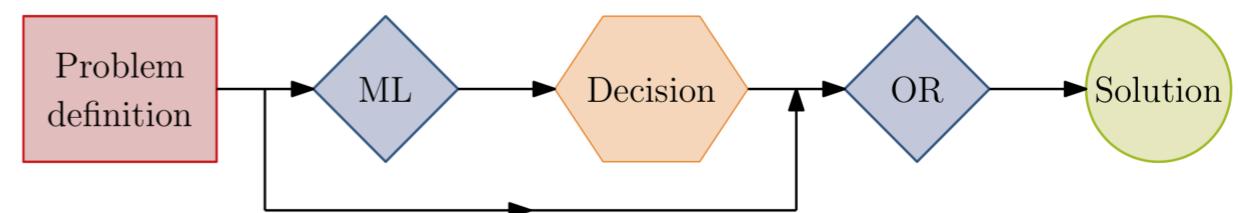
- **End to End Learning**

- ▶ Build a machine learning model to directly output solutions.



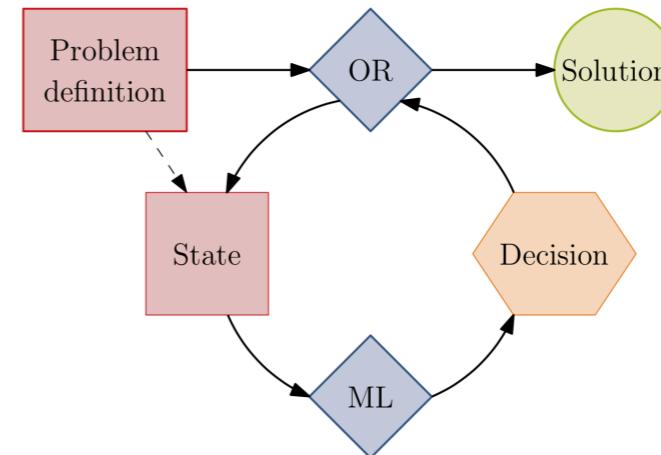
- **Learning Properties**

- ▶ Build a machine learning model to parametrize your OR algorithm.



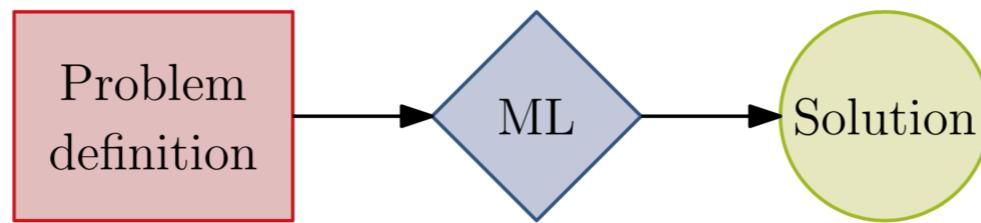
- **Learning Repeated Decisions**

- ▶ The machine learning model is used repeatedly by the OR algorithm to assist it in its solving process.



# Methodology - Paradigms

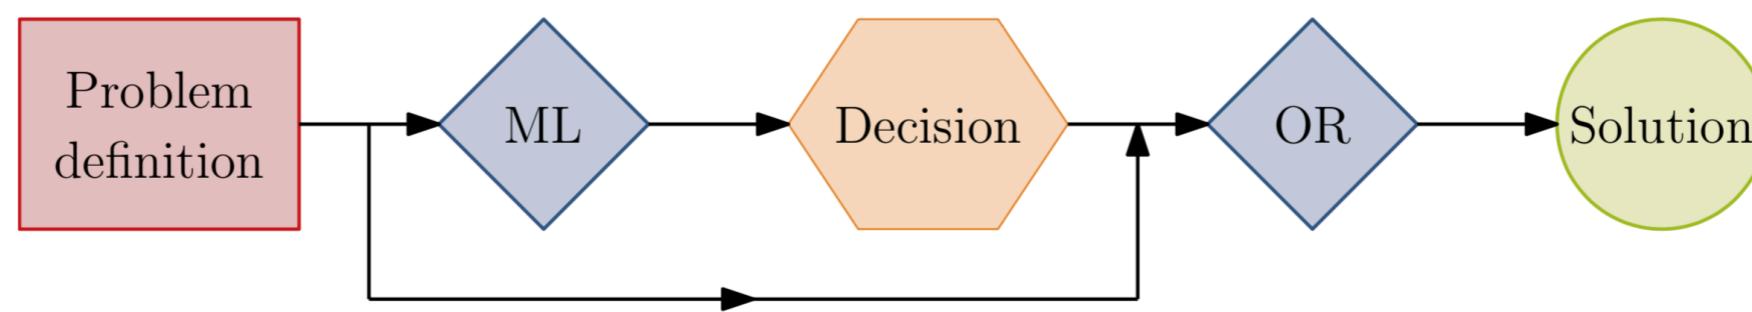
- **End to End Learning**



- Learning TSP solutions  
[Bello et al. 2017][Kool and Welling 2018][Emami and Ranka 2018]  
[Vinyals et al. 2015][Nowak et al. 2017]
- Predict aggregated solutions to MILP under partial information  
[Larsen et al. 2018]
- Approximate obj value to SDP (for cut selection)  
[Baltean-Lugojan et al. 2018]

# Methodology - Paradigms

- **Learning Properties**

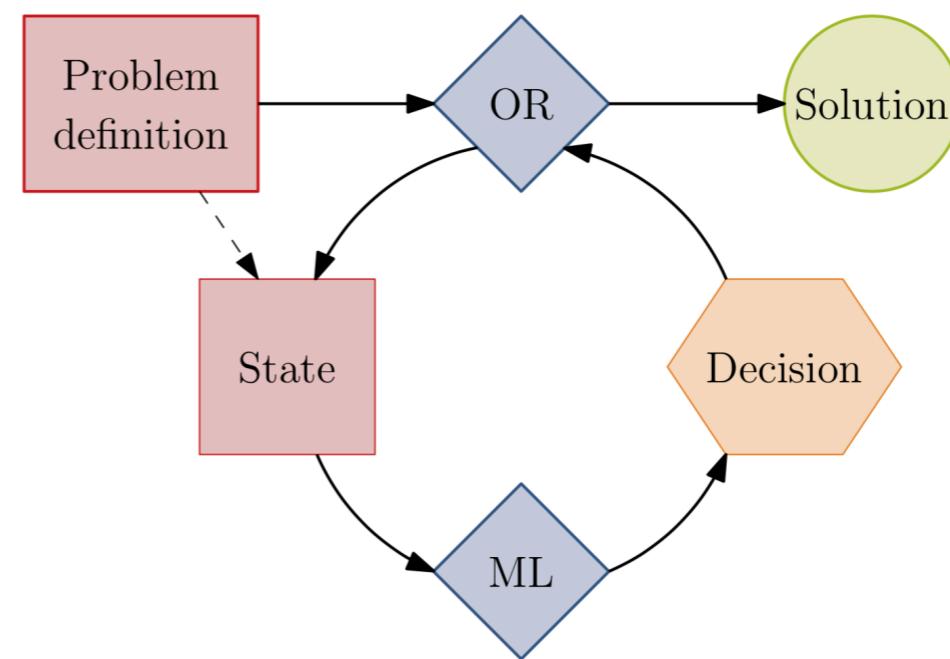


- Use a decomposition method  
*[Kruber et al. 2017]*
- Linearize an MIQP  
*[Bonami et al. 2018]*
- Provide initial cancer treatment plans to inverse optimization  
*[Mahmood et al. 2018]*

# Methodology - Paradigms

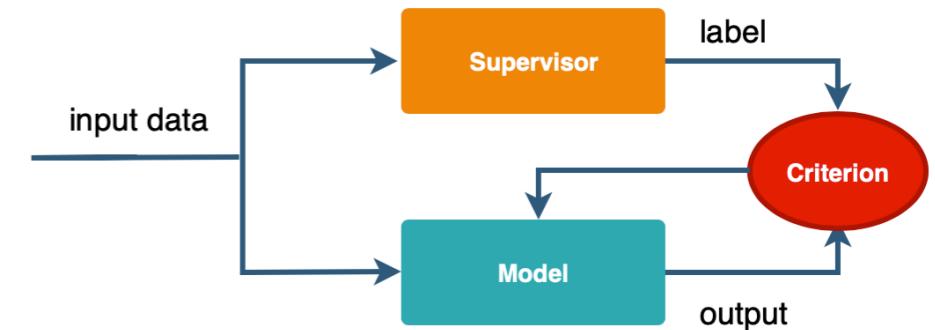
- **Learning Repeated Decisions**

- Learning where to run heuristics in B&B  
*[Khalil et al. 2017b]*
- Learning to branch  
*[Lodi and Zarpellon 2017]* (survey)
- Learning gradient descent  
e.g. *[Andrychowicz et al. 2016]*
- Learning cutting plane selection  
*[Baltean-Lugojan et al. 2018]*

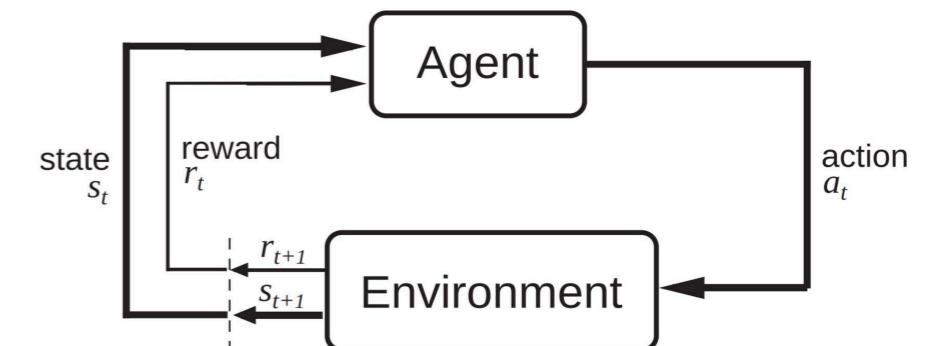


# Methodology - Learning Mechanisms

- By label/demonstration:
  - approximating branching heuristics [Khalil et al. 2016]
  - predicting outcome of cuts [Baltean-Lugojan et al. 2018]
  - approximating tactical solutions [Larsen et al. (2019)]
- By experience
  - learning to solve TSP [Bello et al. 2016, Kool et al. 2018]
  - learning to solve problems on graphs [Khalil et al. 2017a]
  - learning to select cuts [Tang et al. 2019]



Supervised Learning



RL framework by Sutton et al., 1998

# Methodology - Representation

## Applications of Deep Neural Networks:

- *Multilayer Perceptron (MLPs)*
  - fixed-sized data [Larsen et al. 2019]
- *Recurrent Neural Networks (RNNs)*
  - sequential data [Vinyals et al. 2015, Bello et al. 2016, ]
- *Graph Neural Networks (GNNs)*
  - graph space data [Khalil et al. 2017a, Kool et al. 2018, Gasse et al. 2019]
- *Convolutional Neural Networks (CNNs)*
  - pixel space data [Silver et al., 2016, 2018, Miki et all. 2019]

# Outline

- ▶ Introduction
- ▶ ML for combinatorial optimization
- ▶ **Applications with GNNs**

## Learn heuristics for graph opt. problems

Khalil et al. "Learning combinatorial optimization algorithms over graphs." *NIPS*. 2017.

# Learn heuristics for graph opt. problems

Khalil et al. "Learning combinatorial optimization algorithms over graphs." NIPS. 2017.

- Problem statement

Given a **graph optimization problem  $G$**  and a **distribution  $\mathcal{D}$**  of problem instances, can we **learn better greedy heuristics** that generalize to unseen instances from  $\mathcal{D}$ ?

**Minimum Vertex Cover**

Insert nodes into cover

**Maximum Cut**

Insert nodes into subset

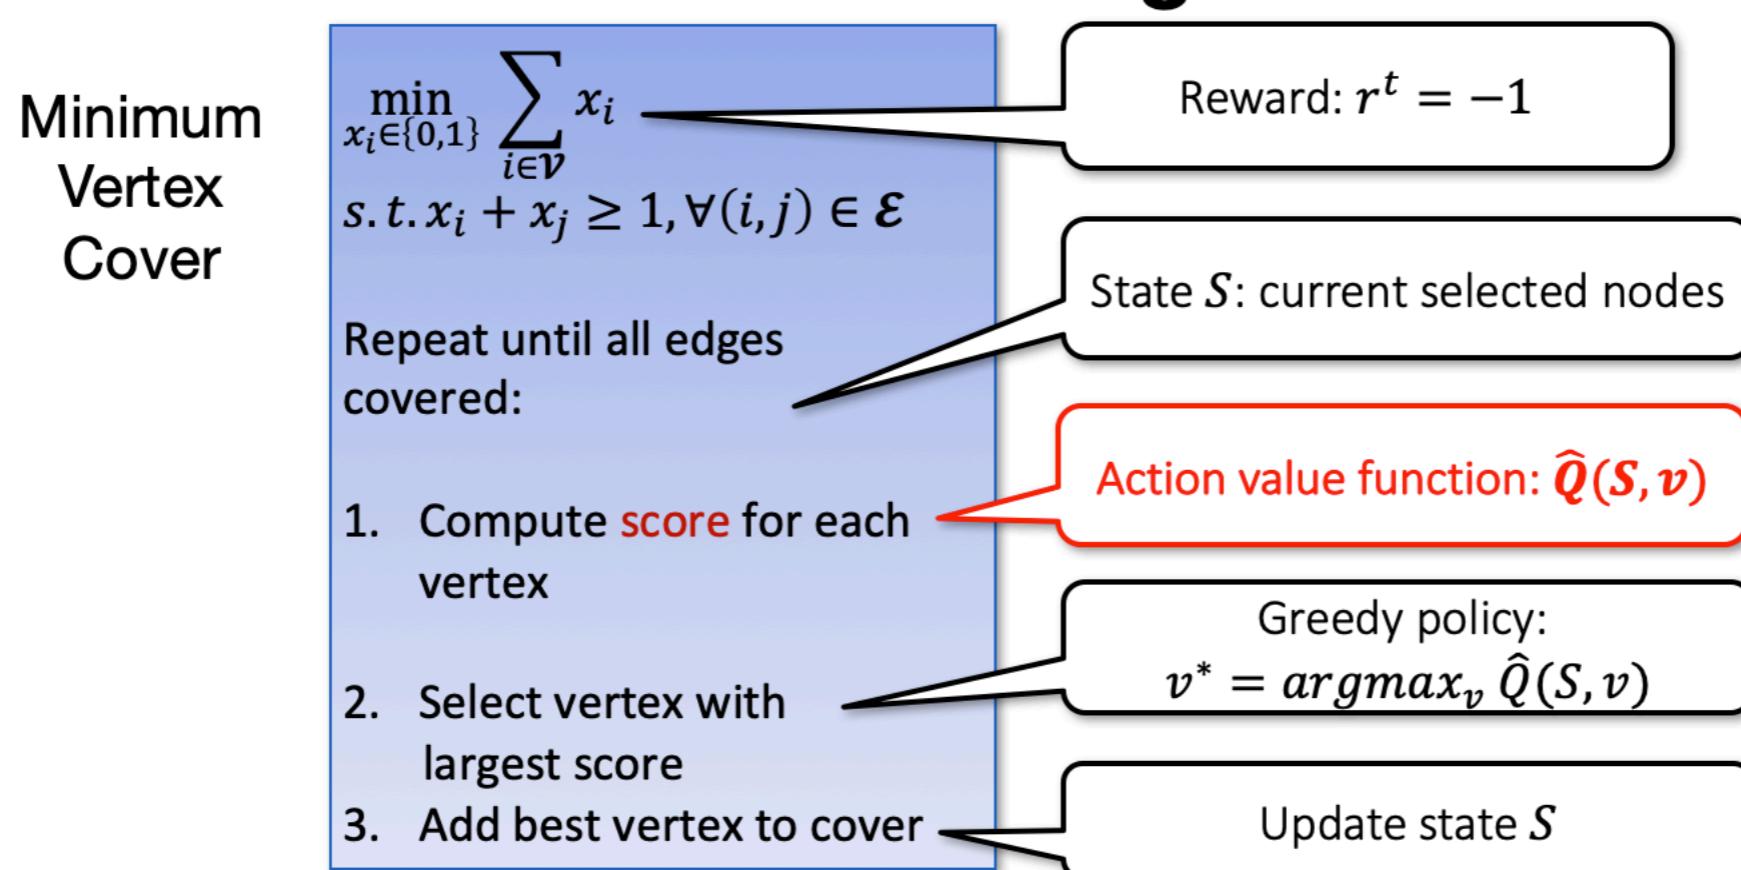
**Traveling Salesman Prob.**

Insert nodes into sub-tour

# Learn heuristics for graph opt. problems

Khalil et al. "Learning combinatorial optimization algorithms over graphs." NIPS. 2017.

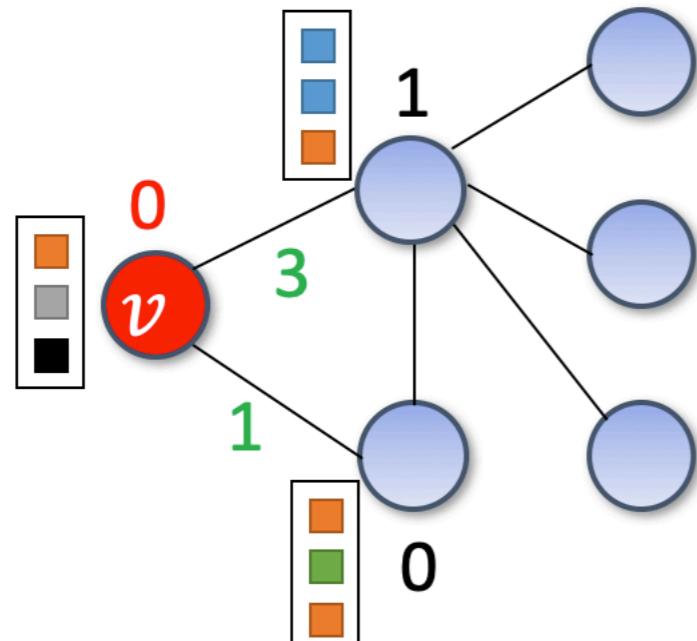
- RL Q-learning framework



# Learn heuristics for graph opt. problems

Khalil et al. "Learning combinatorial optimization algorithms over graphs." NIPS. 2017.

- Graph embedding (structure2vector)



**Non-linearity:**  $\text{relu}(x) = \max(0, x)$

Repeat embedding  $T$  times:

Updating feature vector

$$\mu_v^{(t+1)} \leftarrow \text{relu}(\theta_1 x_v + \frac{\theta_2 \sum_{u \in \mathcal{N}(v)} \mu_u^{(t)} + \theta_3 \sum_{u \in \mathcal{N}(v)} \text{relu}(\theta_4 w(v, u))}{\text{Neighbors' edge weights}})$$

**Θ: model parameters**

# Learn heuristics for graph opt. problems

Khalil et al. "Learning combinatorial optimization algorithms over graphs." NIPS. 2017.

- Results on realistic instances

Table 3: Realistic data experiments, results summary. Values are average approximation ratios.

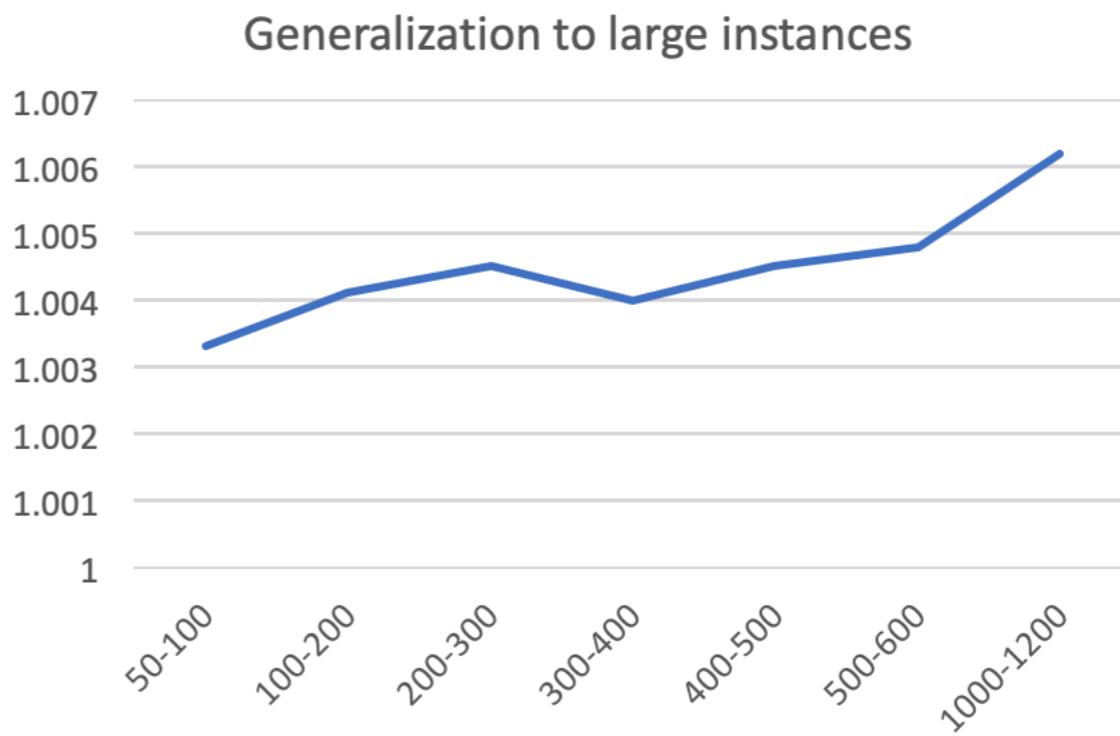
<b>Problem</b>	<b>Dataset</b>	<b>S2V-DQN</b>	<b>Best Competitor</b>	<b>2<sup>nd</sup> Best Competitor</b>
MVC	MemeTracker	<b>1.0021</b>	1.2220 (MVCAprox-Greedy)	1.4080 (MVCAprox)
MAXCUT	Physics	<b>1.0223</b>	1.2825 (MaxcutApprox)	1.8996 (SDP)
TSP	TSPLIB	<b>1.0475</b>	1.0947 (2-opt)	1.1771 (Cheapest)

# Learn heuristics for graph opt. problems

Khalil et al. "Learning combinatorial optimization algorithms over graphs." NIPS. 2017.

- Generalization to larger instances

- Train on small graphs with 50-100 nodes
- Generalize to not only graphs from same distribution
- But also larger graphs
- Approximation ratio < 1.007



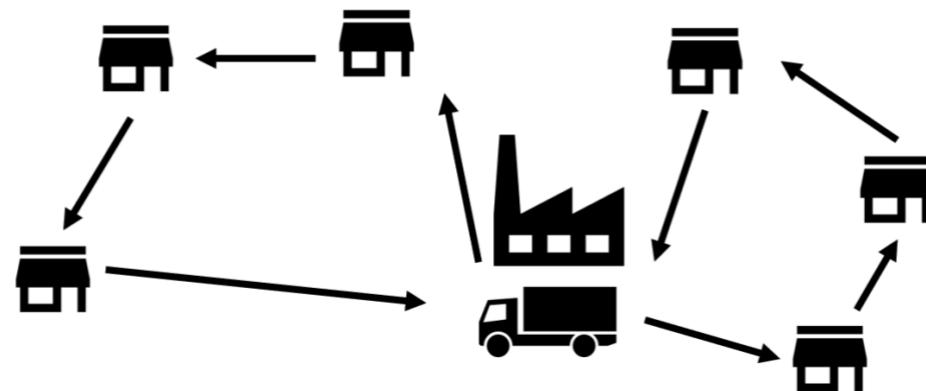
# Learn heuristics for routing problems

Kool et all. "Attention, learn to solve routing problems!." ICLR. 2019.

# Learn heuristics for routing problems

Kool et all. "Attention, learn to solve routing problems!." ICLR. 2019.

- Problem statement



- ▶ Simplest form of the routing problem
  - One capacitated vehicle
  - Multiple customers
  - Vehicle must return to the depot to refill
  - Objective: Minimize total distance

# Previous work: Neural Combinatorial Optimization with RL

Bello et al. "Neural combinatorial optimization with reinforcement learning."(2016).

- Policy model: Pointer Network

$$P_{\theta}(solution|problem)$$

- RL learning:
  - learn from experience by exploring the solution space, no label needed
  - cost: the tour length of the solution
  - training : adjust  $\theta$  with RL to minimize the objective(the expected cost of solutions)

$$J(\theta|problem) = E_{solution \sim P_{\theta^*}(\cdot|problem)}[cost(solution|problem)]$$

$$\theta^* = \operatorname{argmax}_{\theta} J(\theta|problem)$$

- Inference: sample solutions from

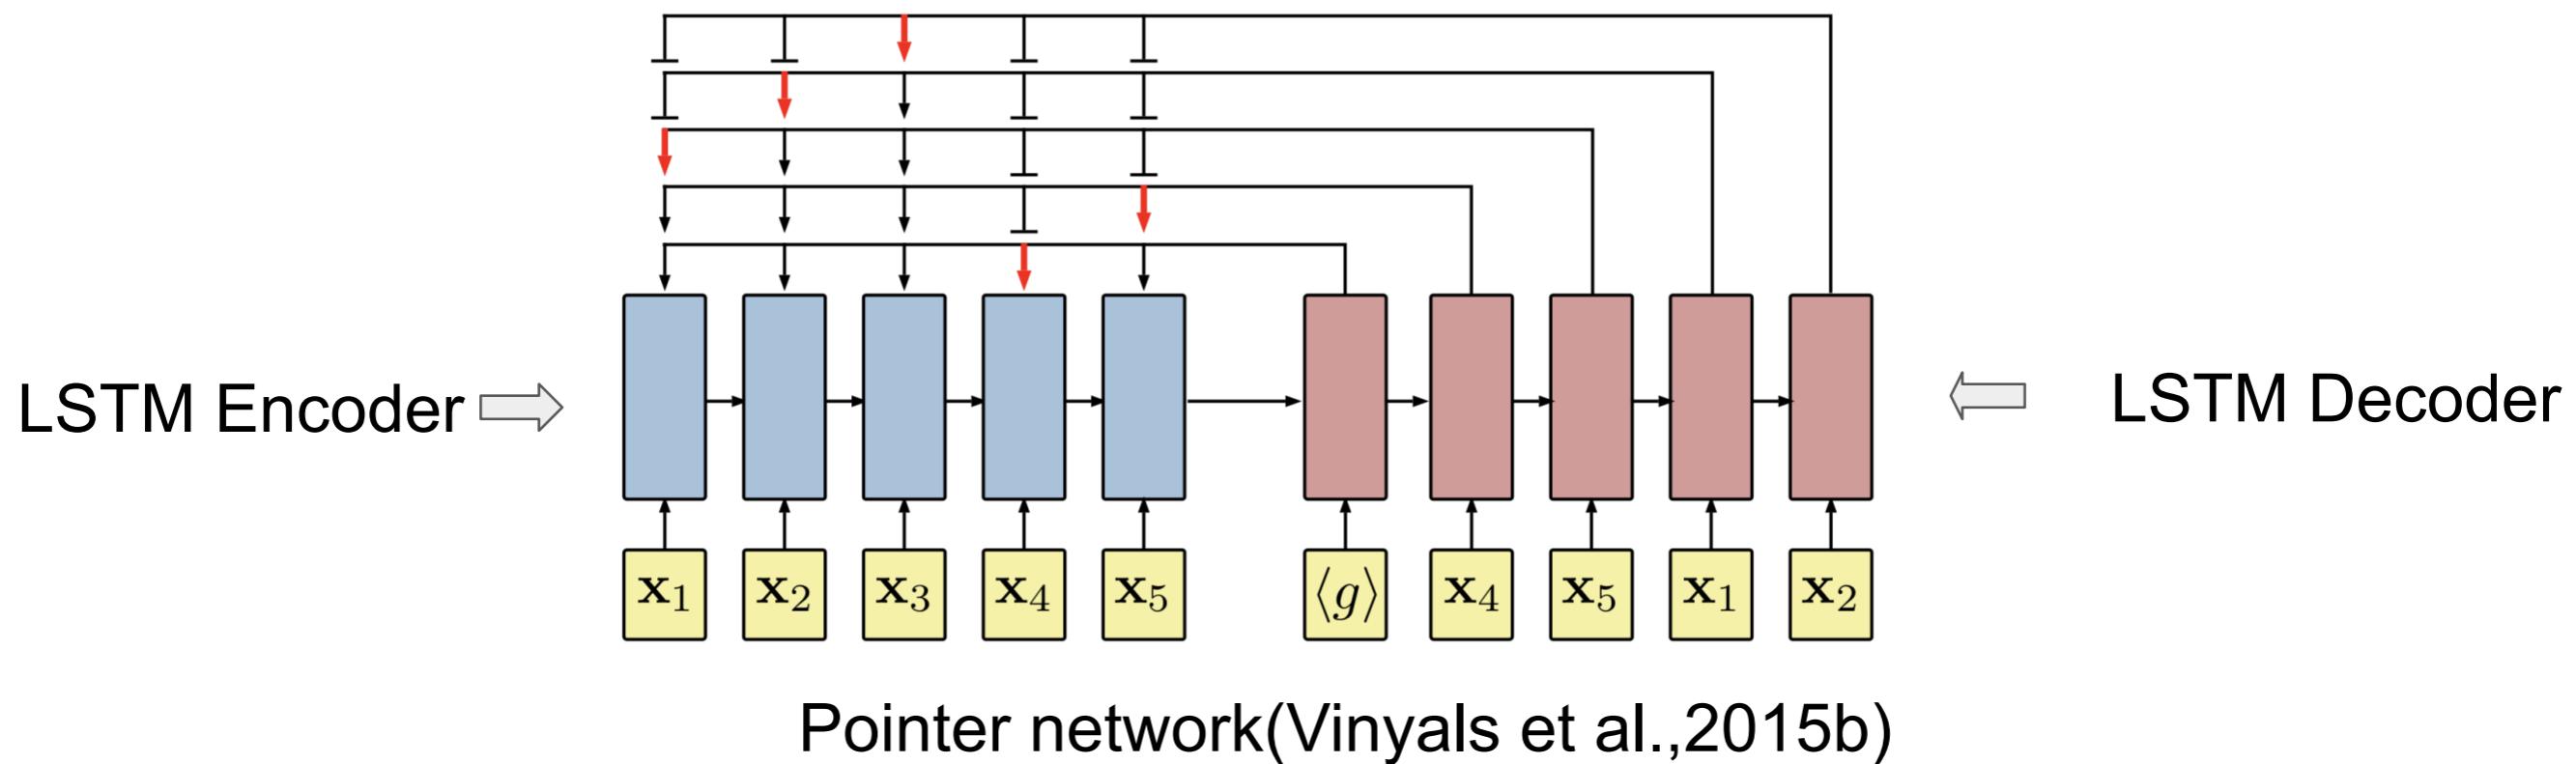
$$solution \sim P_{\theta^*}(\cdot|problem)$$

# Previous work: Neural Combinatorial Optimization with RL

Bello et al. "Neural combinatorial optimization with reinforcement learning."(2016).

- Policy model: Pointer Network

$$P_{\theta}(\text{solution}|\text{problem})$$



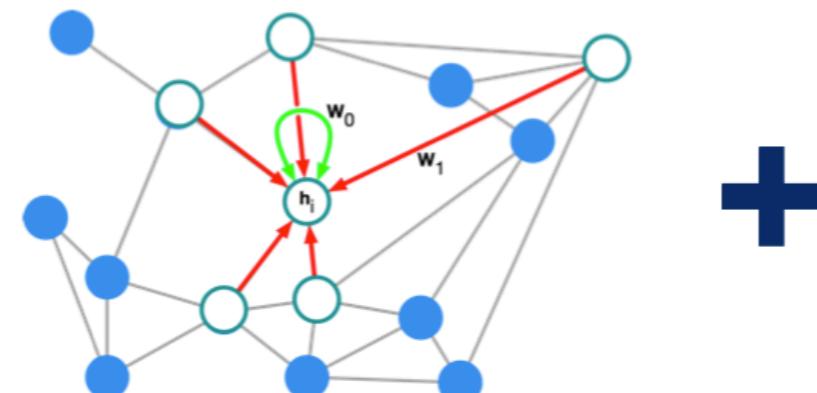
# Learn heuristics for routing problems

Kool et all. "Attention, learn to solve routing problems!." ICLR. 2019.

- Policy model: Pointer Network

$$P_{\theta}(\text{solution}|\text{problem})$$

## Graph convolutions



---

### Attention Is All You Need

---

Ashish Vaswani\*  
Google Brain  
avaswani@google.com      Noam Shazeer\*  
Google Brain  
noam@google.com      Niki Parmar\*  
Google Research  
nikip@google.com      Jakob Uszkoreit\*  
Google Research  
usz@google.com

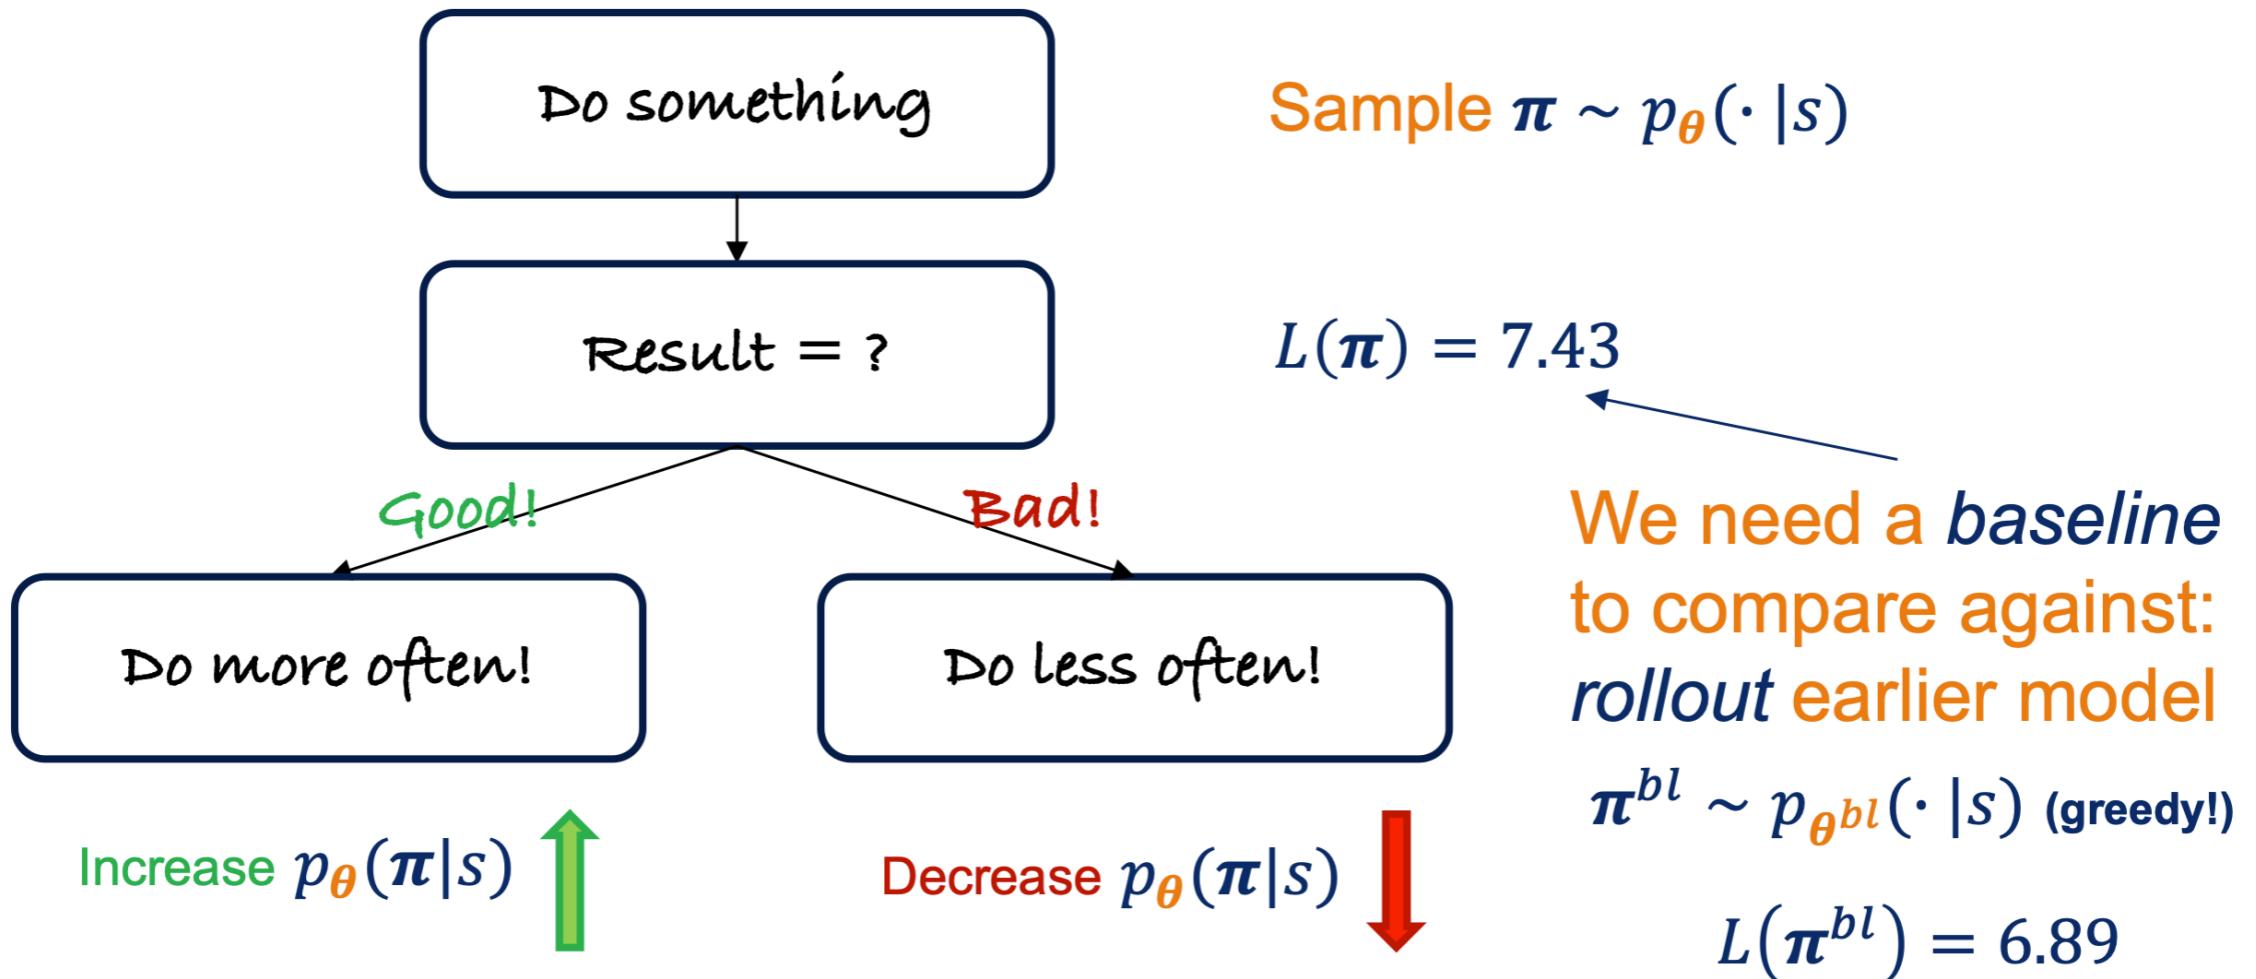
Llion Jones\*  
Google Research  
llion@google.com      Aidan N. Gomez\* †  
University of Toronto  
aidan@cs.toronto.edu      Lukasz Kaiser\*  
Google Brain  
lukaszkaiser@google.com

Illia Polosukhin\* ‡  
illia.polosukhin@gmail.com

# Learn heuristics for routing problems

Kool et all. "Attention, learn to solve routing problems!." ICLR. 2019.

- Learning mechanism: Reinforce with rollout baseline



# Learn heuristics for routing problems

Kool et all. "Attention, learn to solve routing problems!." ICLR. 2019.

- Results: GNNs-Attention model + rollout baseline
  - Improves over classical heuristics!
  - Improves over prior learned heuristics!
    - Attention Model improves
    - Rollout helps significantly
  - Gets close to single-purpose SOTA (20 to 100 nodes)!
    - TSP 0.34% to 4.53% (greedy)
    - TSP 0.08% to 2.26% (best of 1280 samples)

## Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Mixed-Integer Linear Program (MILP)
- 

$$\begin{aligned} & \arg \min_{\mathbf{x}} \quad \mathbf{c}^T \mathbf{x} \\ \text{subject to} \quad & \mathbf{A} \mathbf{x} \leq \mathbf{b}, \\ & \mathbf{l} \leq \mathbf{x} \leq \mathbf{u}, \\ & \mathbf{x} \in \mathbb{Z}^p \times \mathbb{R}^{n-p}. \end{aligned}$$

- ▶  $\mathbf{c} \in \mathbb{R}^n$  the objective coefficients
- ▶  $\mathbf{A} \in \mathbb{R}^{m \times n}$  the constraint coefficient matrix
- ▶  $\mathbf{b} \in \mathbb{R}^m$  the constraint right-hand-sides
- ▶  $\mathbf{l}, \mathbf{u} \in \mathbb{R}^n$  the lower and upper variable bounds
- ▶  $p \leq n$  integer variables

# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Linear Relaxation

$$\begin{aligned} & \arg \min_{\mathbf{x}} \quad \mathbf{c}^T \mathbf{x} \\ \text{subject to} \quad & \mathbf{A}\mathbf{x} \leq \mathbf{b}, \\ & \mathbf{l} \leq \mathbf{x} \leq \mathbf{u}, \\ & \mathbf{x} \in \mathbb{R}^n. \end{aligned}$$

Convex problem, efficient algorithms (e.g., simplex).

- ▶  $\mathbf{x}^* \in \mathbb{Z}^p \times \mathbb{R}^{n-p}$  (lucky)  $\rightarrow$  solution to the original MILP
- ▶  $\mathbf{x}^* \notin \mathbb{Z}^p \times \mathbb{R}^{n-p} \rightarrow$  lower bound to the original MILP

# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Linear Relaxation

$$\begin{aligned} & \arg \min_{\mathbf{x}} \quad \mathbf{c}^T \mathbf{x} \\ \text{subject to} \quad & \mathbf{A}\mathbf{x} \leq \mathbf{b}, \\ & \mathbf{l} \leq \mathbf{x} \leq \mathbf{u}, \\ & \mathbf{x} \in \mathbb{R}^n. \end{aligned}$$

Convex problem, efficient algorithms (e.g., simplex).

- ▶  $\mathbf{x}^* \in \mathbb{Z}^p \times \mathbb{R}^{n-p}$  (lucky)  $\rightarrow$  solution to the original MILP
- ▶  $\mathbf{x}^* \notin \mathbb{Z}^p \times \mathbb{R}^{n-p} \rightarrow$  lower bound to the original MILP

# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Branch & Bound

.

Split the LP recursively over a non-integral variable, i.e.  $\exists i \leq p \mid x_i^* \notin \mathbb{Z}$

$$x_i \leq \lfloor x_i^* \rfloor \quad \vee \quad x_i \geq \lceil x_i^* \rceil.$$

**Lower bound (L)**: minimal among leaf nodes.

**Upper bound (U)**: minimal among leaf nodes with integral solution.

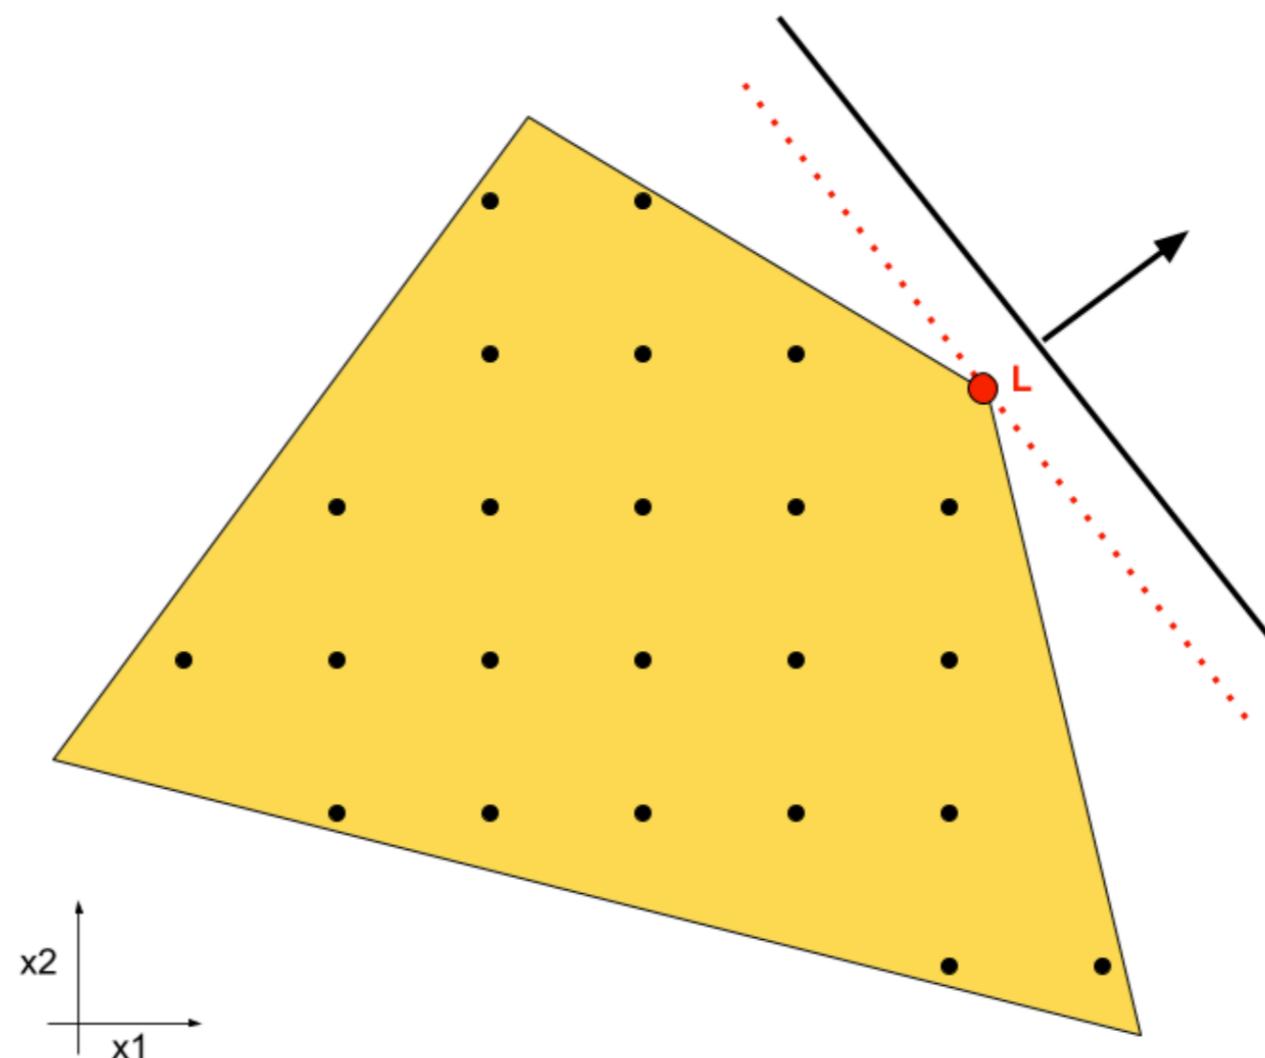
Stopping criterion:

- ▶ **L = U** (optimality certificate)
- ▶ **L =  $\infty$**  (infeasibility certificate)
- ▶ **L - U < threshold** (early stopping)

# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

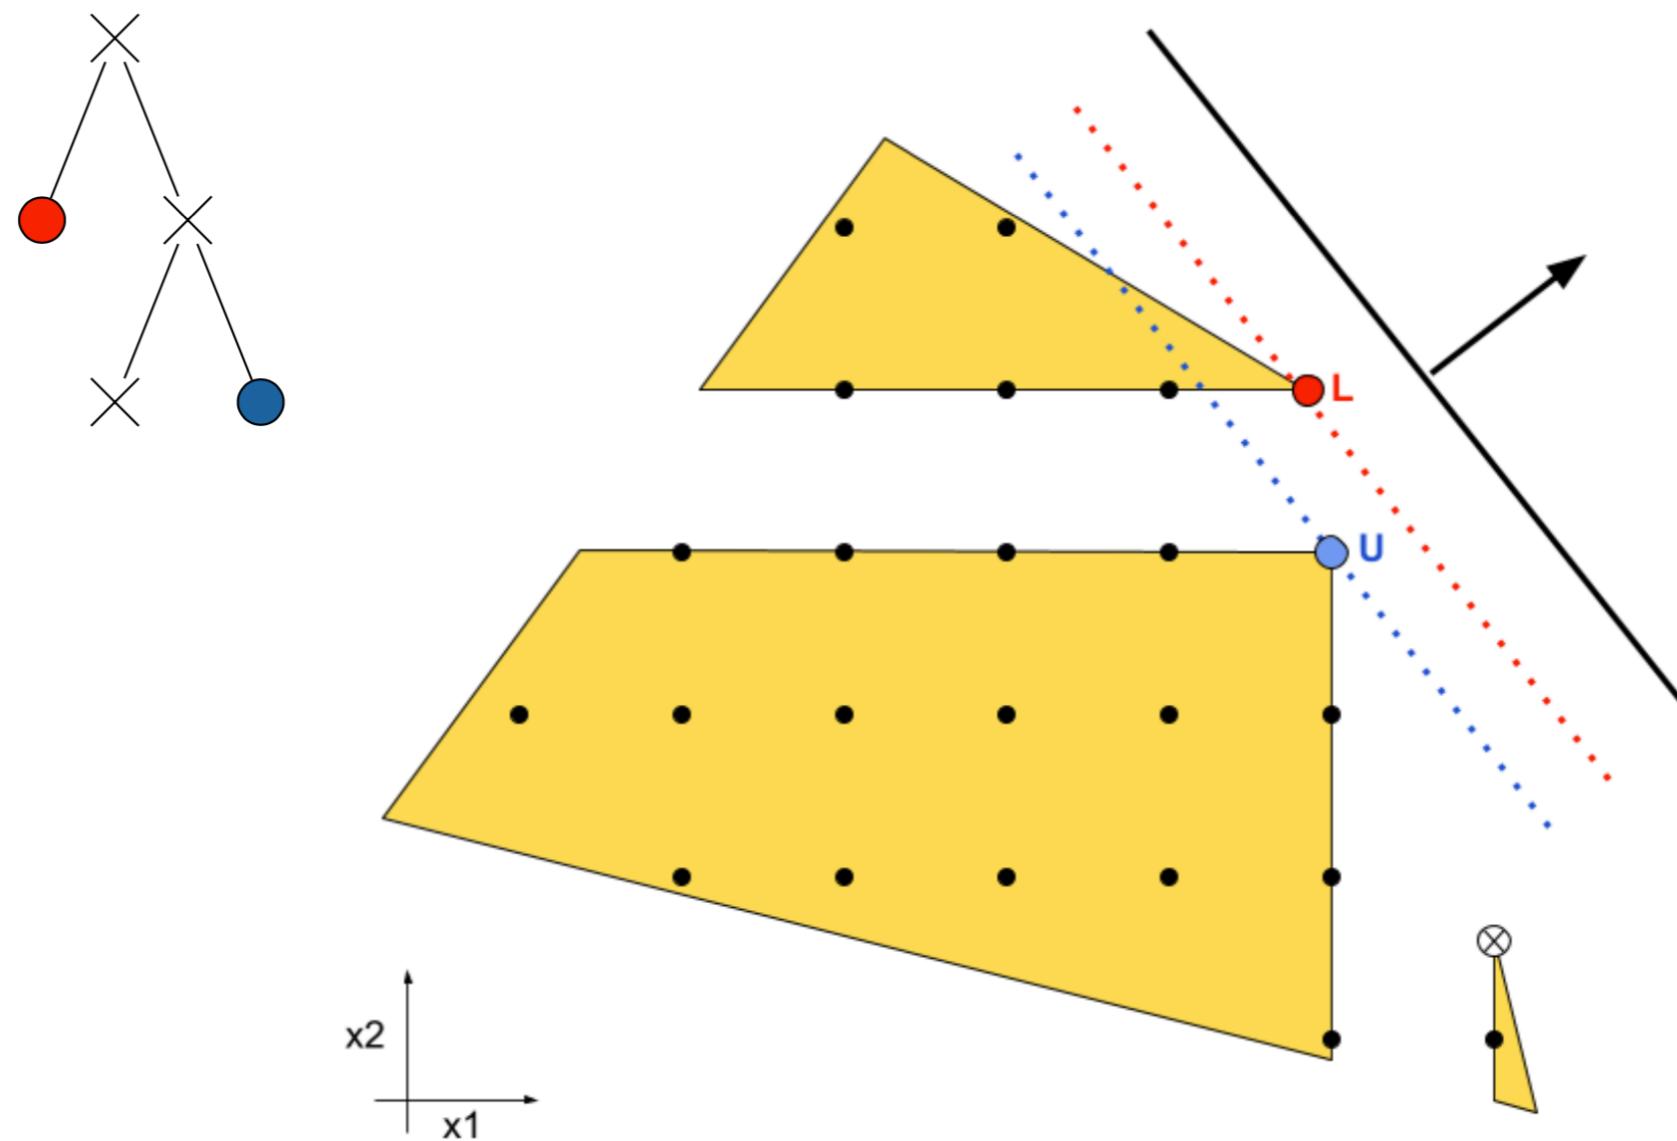
- Branch & Bound
- 



# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

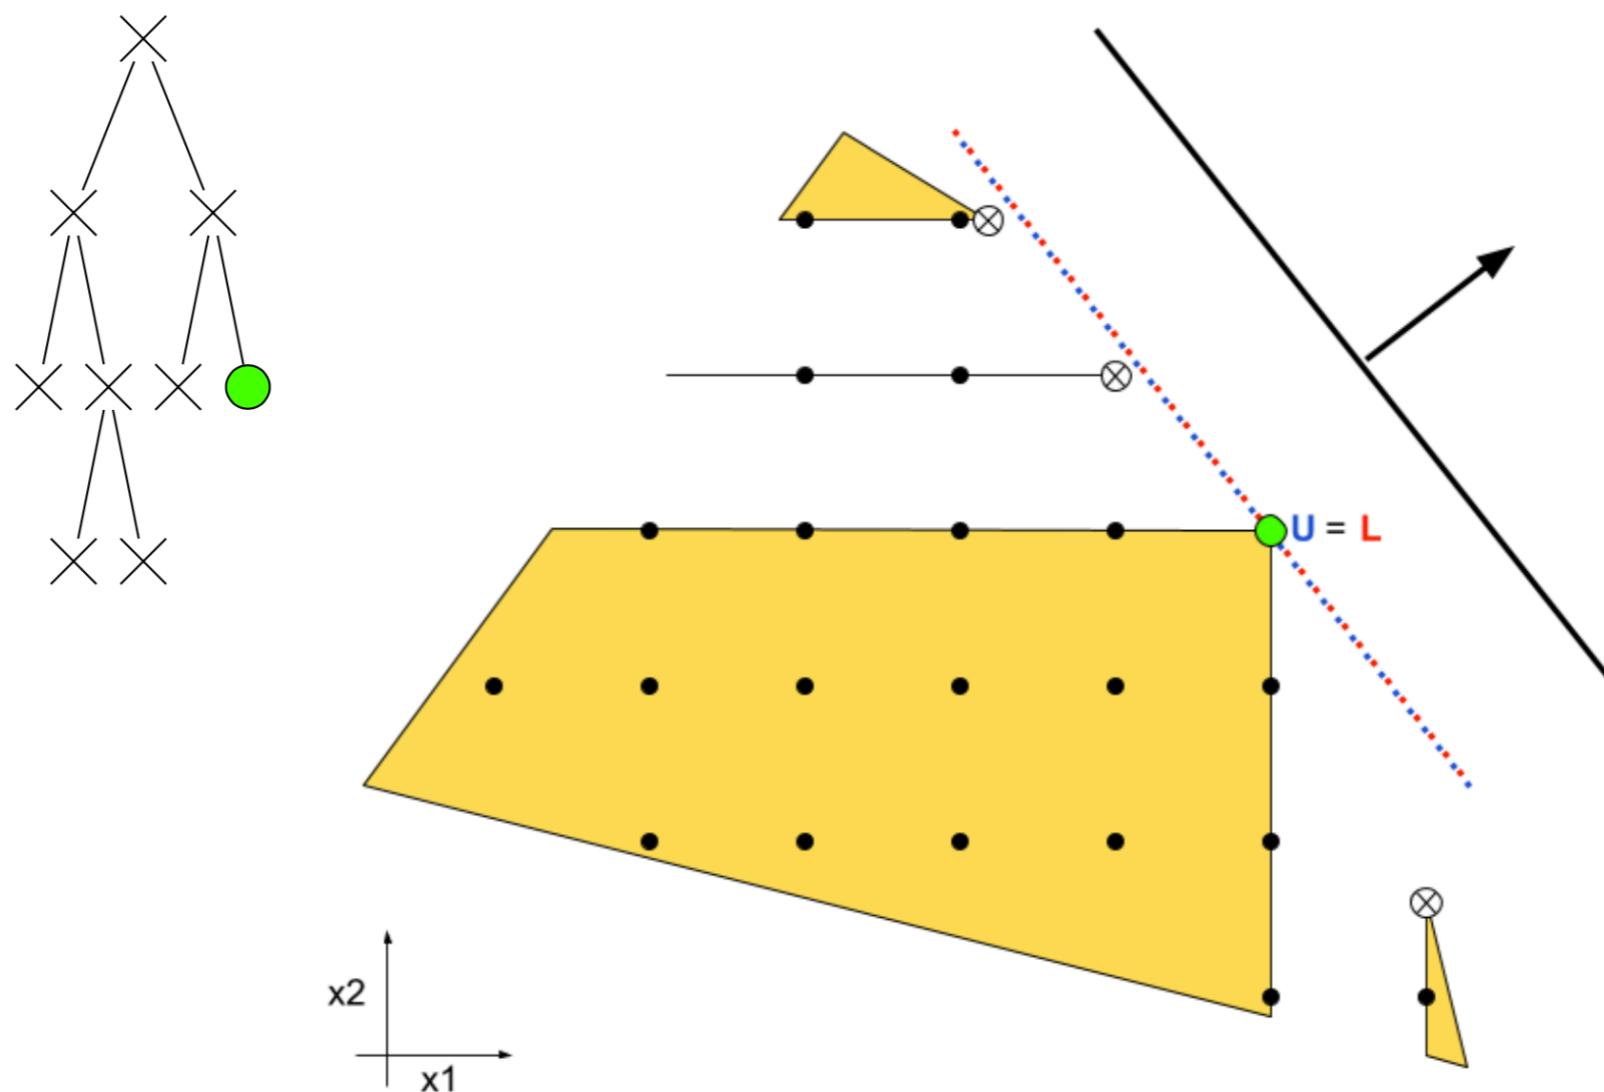
- Branch & Bound
- 



# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Branch & Bound
- 



# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Expert branching rules : state-of-the-art

Strong branching: one-step forward looking (greedy)

- ▶ solve both LPs for each candidate variable
- ▶ pick the variable resulting in tightest relaxation
- + small trees
- computationally expensive

Pseudo-cost: backward looking

- ▶ keep track of tightenings in past branchings
- ▶ pick the most promising variable
- + very fast, almost no computations
- cold start

Reliability pseudo-cost: best of both worlds

- ▶ compute SB scores at the beginning
- ▶ gradually switches to pseudo-cost (+ other heuristics)
- + best overall solving time trade-off (on MIPLIB)

# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Imitation learning

Full Strong Branching (FSB): good branching rule, but expensive.  
Can we learn a fast, good-enough approximation ?

## Behavioural cloning

- ▶ collect  $\mathcal{D} = \{(s, a^*), \dots\}$  from the expert agent (FSB)
- ▶ estimate  $\pi^*(a | s)$  from  $\mathcal{D}$
- + no reward function, supervised learning, well-behaved
- will never surpass the expert...

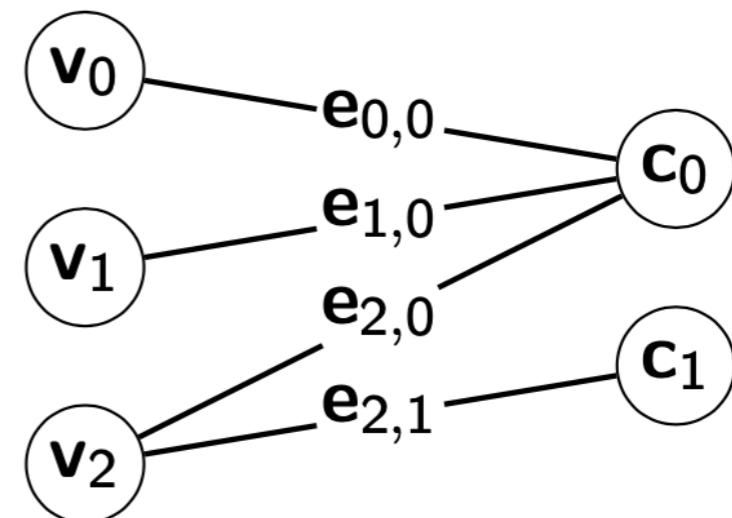
# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- GNNs model

Natural representation : variable / constraint bipartite graph

$$\begin{aligned} \arg \min_{\mathbf{x}} \quad & \mathbf{c}^T \mathbf{x} \\ \text{subject to} \quad & \mathbf{A}\mathbf{x} \leq \mathbf{b}, \\ & \mathbf{l} \leq \mathbf{x} \leq \mathbf{u}, \\ & \mathbf{x} \in \mathbb{Z}^p \times \mathbb{R}^{n-p}. \end{aligned}$$



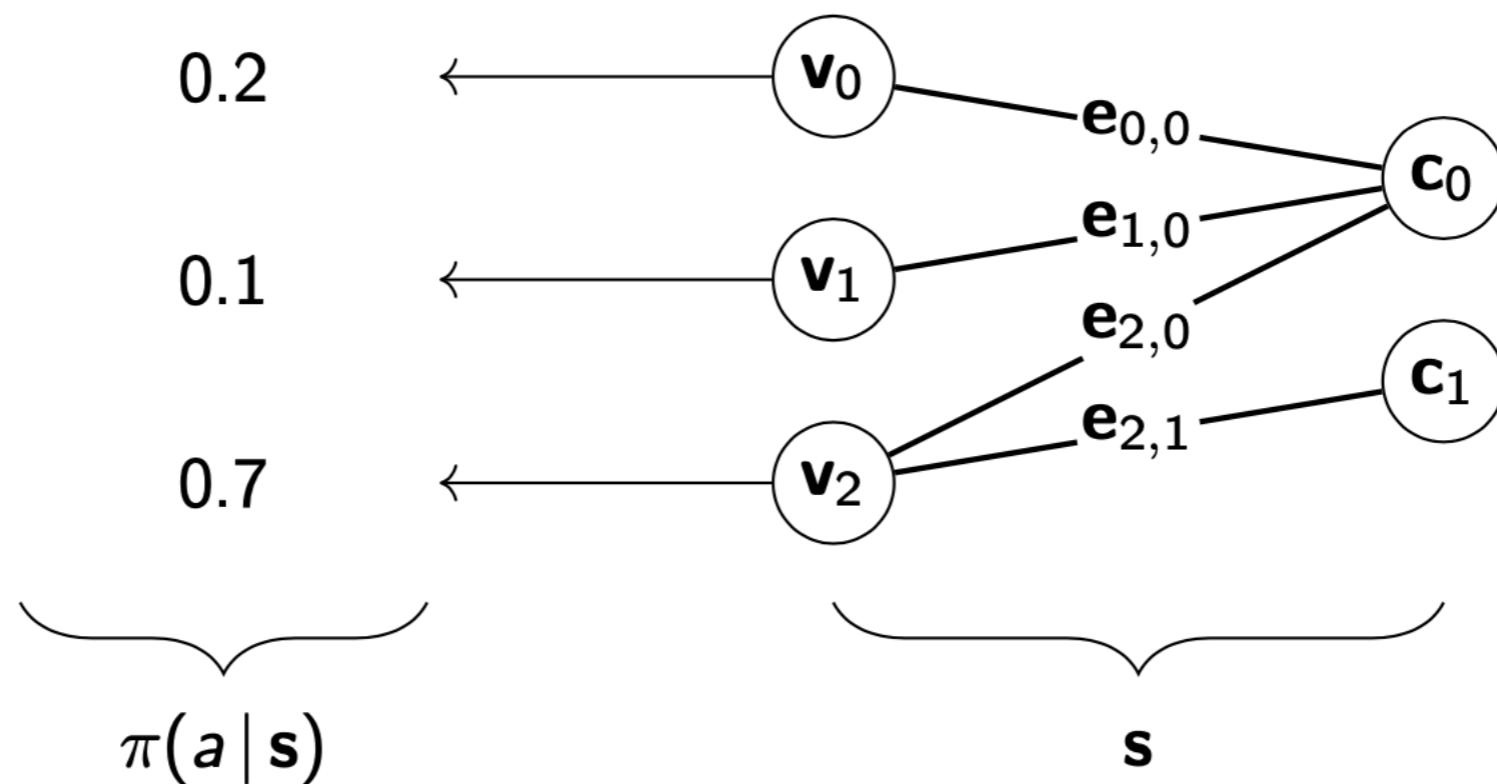
- ▶  $v_i$ : variable features (type, coef., bounds, LP solution...)
- ▶  $c_j$ : constraint features (right-hand-side, LP slack...)
- ▶  $e_{i,j}$ : non-zero coefficients in  $\mathbf{A}$

# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- GNNs model

Neighbourhood-based updates:  $\mathbf{v}_i \leftarrow \sum_{j \in \mathcal{N}_i} \mathbf{f}_\theta(\mathbf{v}_i, \mathbf{e}_{i,j}, \mathbf{c}_j)$



# Learn to branch

Gasse, et al. "Exact combinatorial optimization with graph convolutional neural networks." Neurips 2019

- Experiment ( minimum set covering )

Model	Time	Easy			Medium			Hard		
		Wins	Nodes	Time	Wins	Nodes	Time	Wins	Nodes	
FSB	17.30	0 / 100	17	411.34	0 / 90	171	3600.00	0 / 0	n	
RPB	8.98	0 / 100	<b>54</b>	60.07	0 / 100	1741	1677.02	4 / 65	47	
XTrees	9.28	0 / 100	187	92.47	0 / 100	2187	2869.21	0 / 35	59	
SVMrank	8.10	1 / 100	165	73.58	0 / 100	1915	2389.92	0 / 47	42	
$\lambda$ -MART	7.19	14 / 100	167	59.98	0 / 100	1925	2165.96	0 / 54	45	
GCNN	<b>6.59</b>	<b>85</b> / 100	134	<b>42.48</b>	<b>100</b> / 100	<b>1450</b>	<b>1489.91</b>	<b>66</b> / 70	<b>29</b>	

3 problem sizes

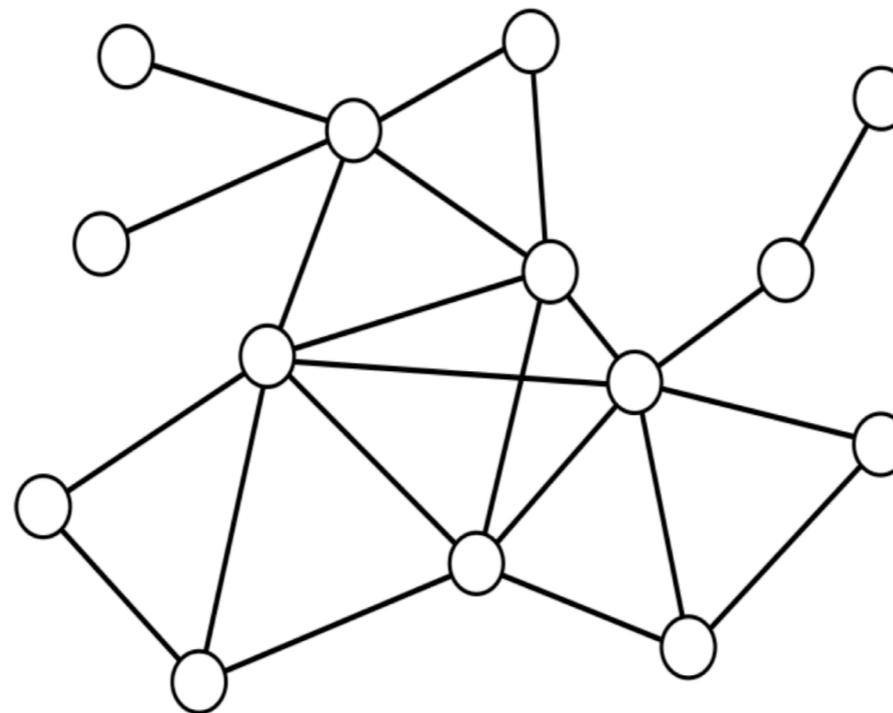
- ▶ 500 rows, 1000 cols (easy), training distribution
- ▶ 1000 rows, 1000 cols (medium)
- ▶ 2000 rows, 1000 cols (hard)

## Learn to build chordal graphs

*Liu et all. "Learning chordal extensions." arXiv preprint arXiv:1910.07600 (2019).*

# Chordal Graphs

A *chordal graph* is an undirected graph with the property that every cycle of length greater than three has a chord (an edge between nonconsecutive vertices in the cycle) [Vandenberghe et al., 2015].

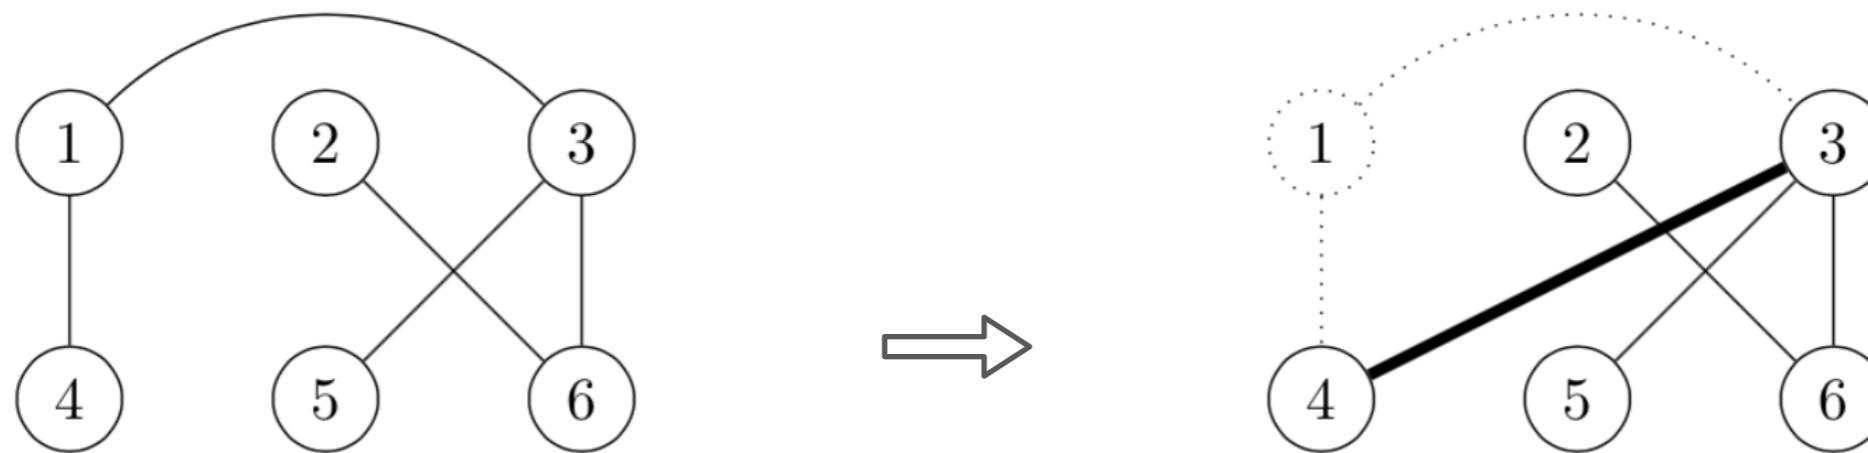


example: a chordal graph

Applications: semidefinite optimization, probabilistic graphical models, dynamic programming.

# Graph Elimination and Chordal Extensions

Any undirected graph can be converted to a *chordal graph* by *graph elimination*. The resulting graph is known as *chordal extension*.



An example of graph elimination. Consider an undirected graph  $G=(V, E)$ , and a vertex  $v \in V$ . The elimination graph is obtained by removing vertex  $v$  and its incident edges from  $G$ , and then adding all edges  $(w, w')$  where  $w$  and  $w'$  were adjacent to  $v$  in  $G$ .

# Motivation

Given a sparse linear system

$$\Phi x = b$$

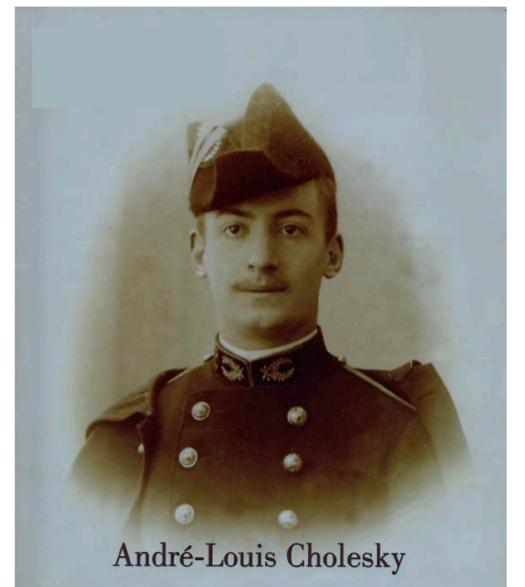
If  $\Phi$  is symmetric and positive definite, we can compute its **Cholesky factorization**

$$\Phi = LL^T$$

Then the original problem is solved by

$$z = L^{-1}b,$$

$$x = L^{-T}z.$$



André-Louis Cholesky

# Motivation

Cholesky factorization:

$$\Phi = LL^T$$

:(  
Cholesky factor L might suffer from significant **fill-in**

$$\begin{pmatrix} \times & \times & \times & \times & \times \\ \times & \times & & & \\ \times & & \times & & \\ \times & & & \times & \\ \times & & & & \times \end{pmatrix} \rightarrow \begin{pmatrix} \times & & & & \\ \times & \times & & & \\ \times & + & \times & & \\ \times & + & + & \times & \\ \times & + & + & + & \times \end{pmatrix}$$

# Motivation

Cholesky factorization:

$$\Phi = LL^T$$

- : ( Cholesky factor L might suffer from significant **fill-in**
- : ) Fill-in can be reduced by re-ordering the rows and columns of the matrix

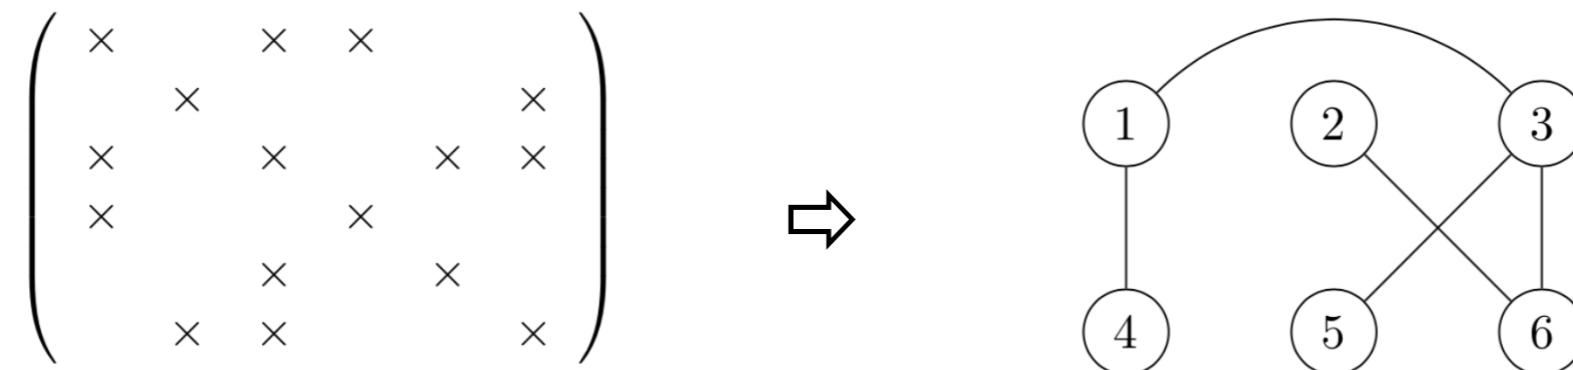
$$\begin{pmatrix} \times & \times & \times & \times & \times \\ \times & \times & & & \\ \times & & \times & & \\ \times & & & \times & \\ \times & & & & \times \end{pmatrix} \rightarrow \begin{pmatrix} \times & & & & \\ \times & \times & & & \\ \times & + & \times & & \\ \times & + & + & \times & \\ \times & + & + & + & \times \end{pmatrix}$$

$$\begin{pmatrix} \times & & & \times & \\ & \times & & \times & \\ & & \times & \times & \\ & & & \times & \times \\ \times & \times & \times & \times & \times \end{pmatrix} \rightarrow \begin{pmatrix} \times & & & & \\ & \times & & & \\ & & \times & & \\ & & & \times & \\ \times & \times & \times & \times & \times \end{pmatrix}$$

# Motivation

For any symmetric matrix  $\Phi$ , its ordered graph is defined as

$$V^\Phi = \{1, \dots, n\},$$
$$E^\Phi = \{(i, j) | \Phi_{i,j} \neq 0, i \neq j\}$$

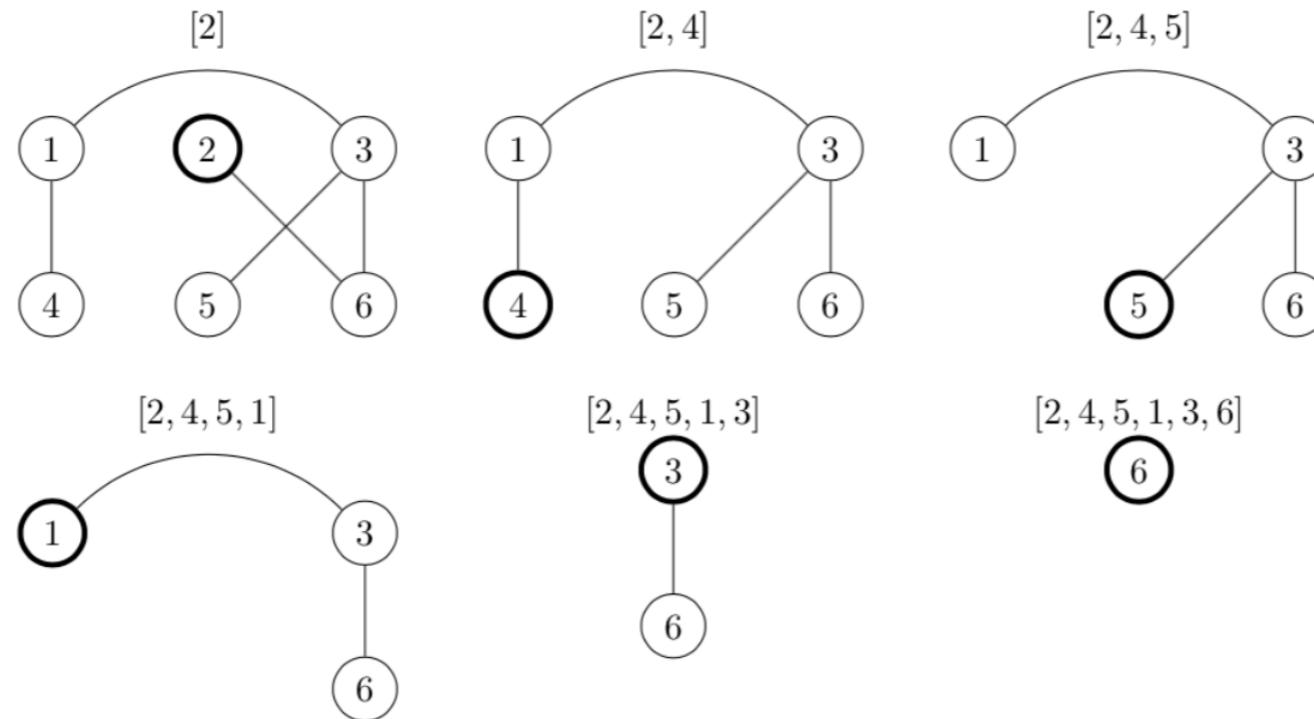


**Objective:** find an ordering that minimizes fill-in. (NP-complete hard)

# Ordering Heuristics

Fill-reducing Heuristics: *minimum degree, nested dissection, and band reduction* [Davis, 2006].

*Minimum degree* : choosing a vertex of minimum degree to be eliminated



Can we leverage ML to learn a better ordering heuristic?

# ML Framework for learning heuristics

- Representation (Policy):

$$P_{\theta}(\text{solution}|\text{problem})$$

- Learning:

- ▶ learn  $\theta$  from a dataset
  - imitation learning
  - reinforcement learning

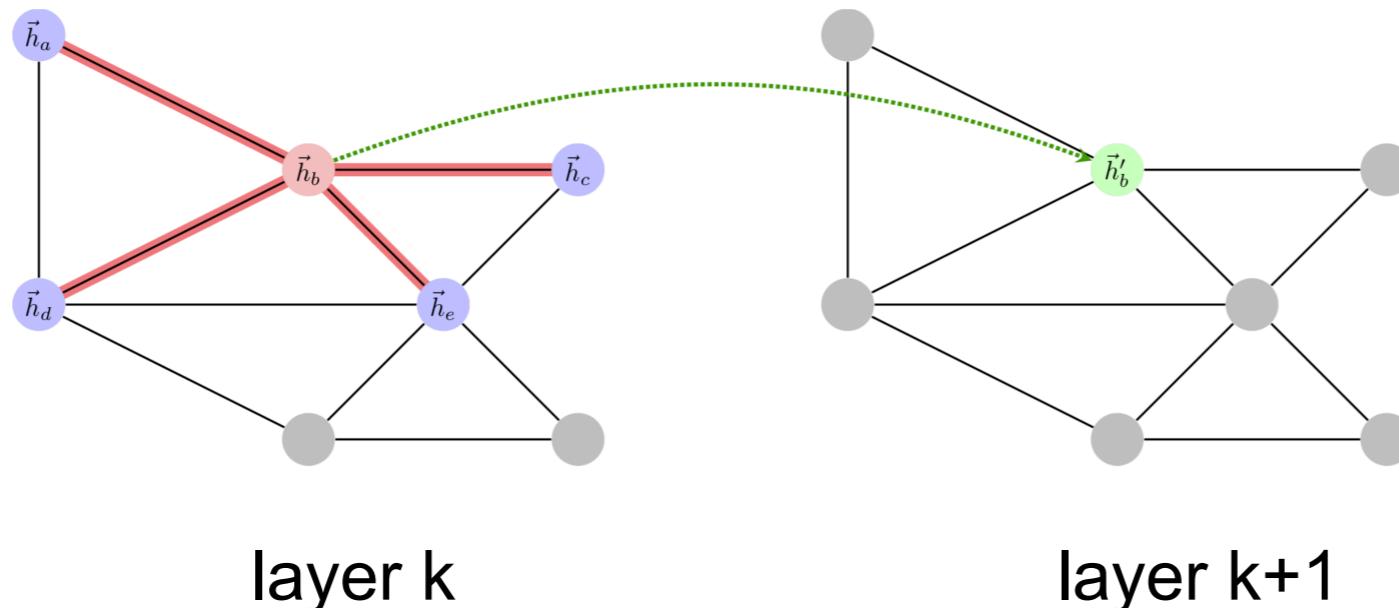
- Inference: sample solutions from

$$\text{solution} \sim P_{\theta^*}(\cdot|\text{problem})$$

# Learning Setting: Representation

- **GNNs Policy Function Approximation:**

$$P_{\theta}(\text{solution}|\text{problem})$$



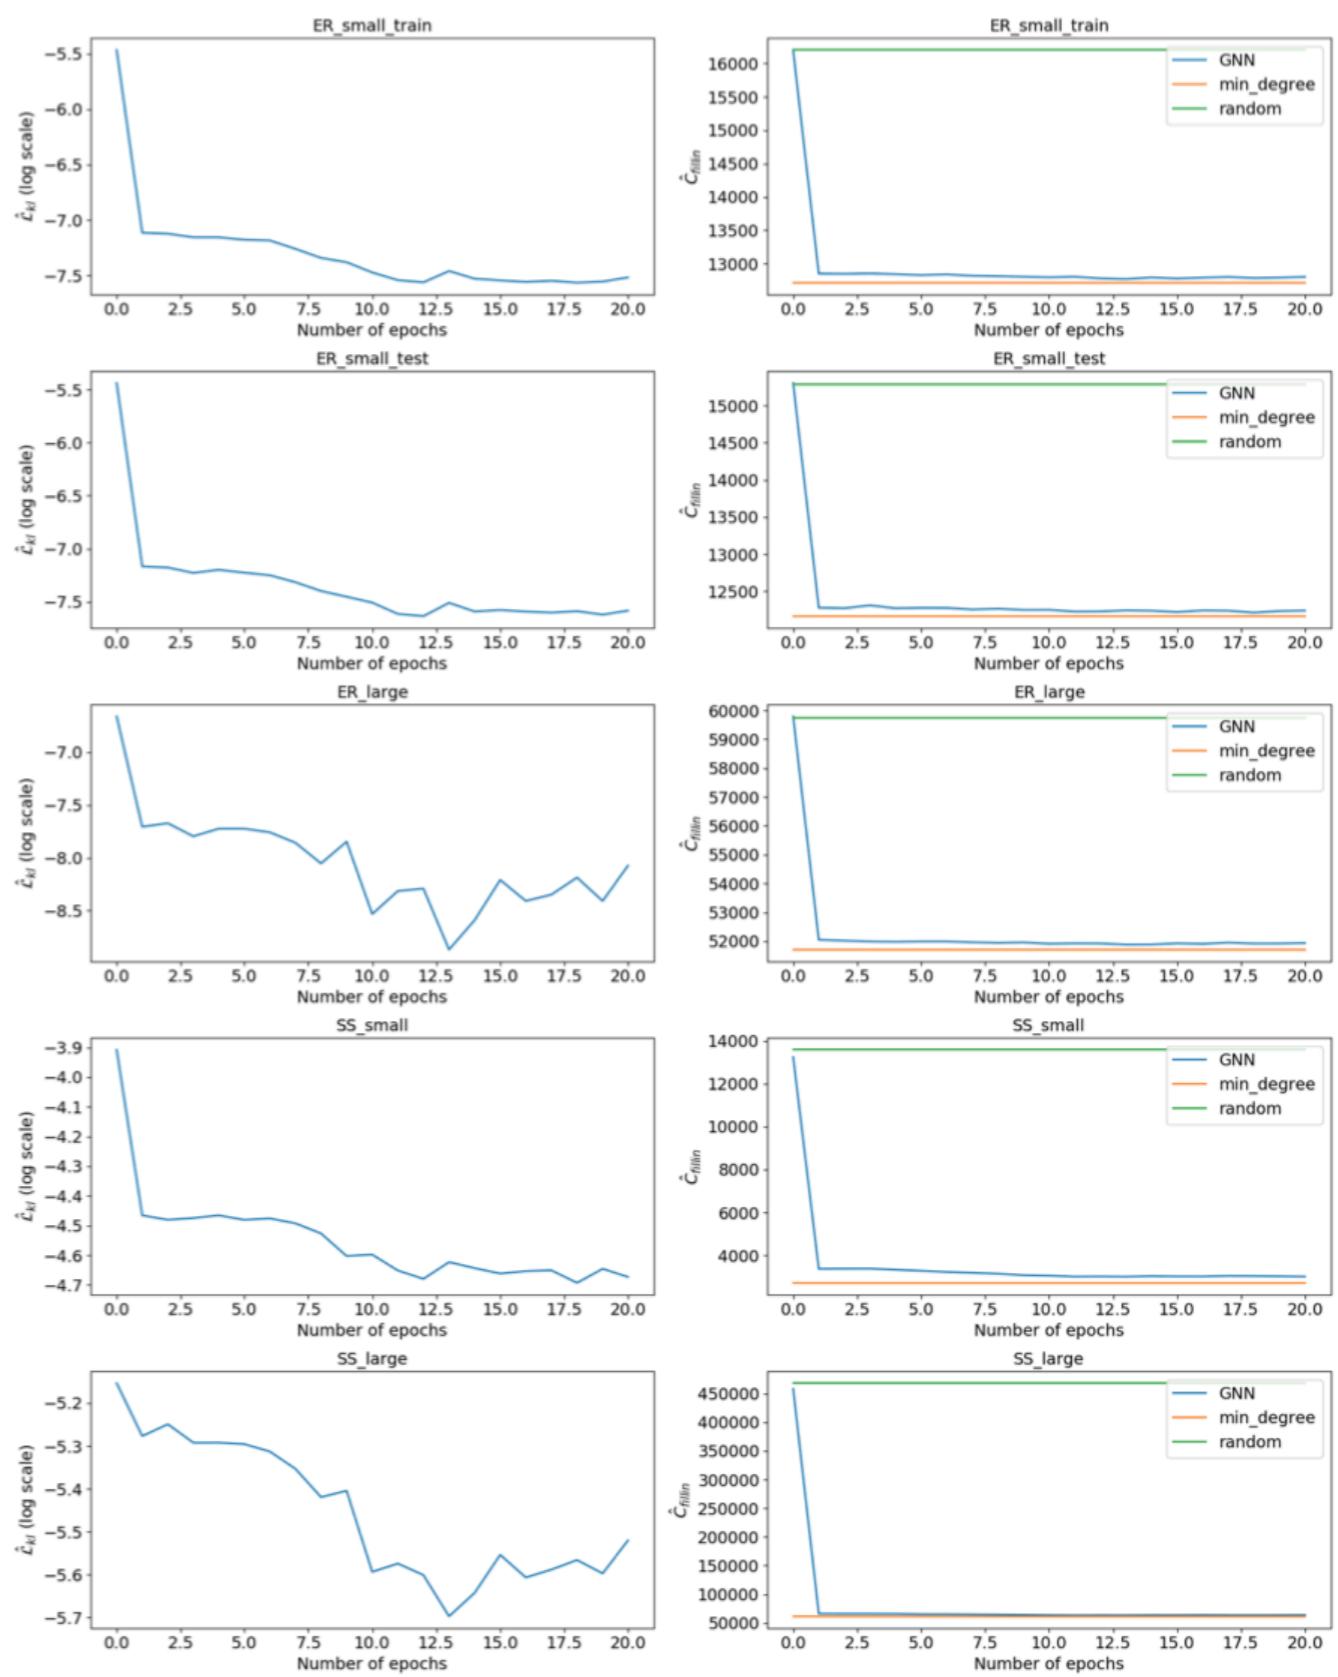
# Learning Setting: Imitation Learning

- **Supervised Learning**
  - learn from dataset(problem, label solution), **so need labels!**
  - training : Adjust  $\theta$  by maximizing the likelihood of the label solutions
$$\theta^* = \operatorname{argmax}_{\theta} P_{\theta}(\text{label solution} | \text{problem})$$
- Idea: can the policy model do the same as the heuristics?
  - Imitation learning:
    - ▶ get labels from an heuristic expert
    - ▶ train the policy model with it iteratively
  - Experimental set-up
    - ▶ 200 Erdos-Renyi random graphs
    - ▶ heuristic experts: minimum degree

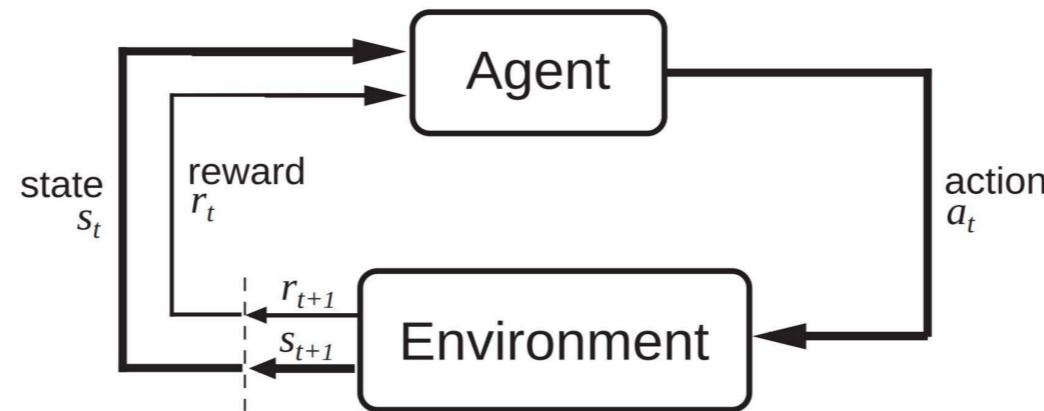
# Preliminary Results:

- Imitation Learning

- Training:
  - ER graphs small(100-200)
- Test:
  - ER graphs small(100-200)
  - ER graphs large(1000-2000)
  - Realistic SS dataset (50-500)
  - Realistic SS dataset (1000-2000)



# Learning Setting: Reinforcement Learning(RL)

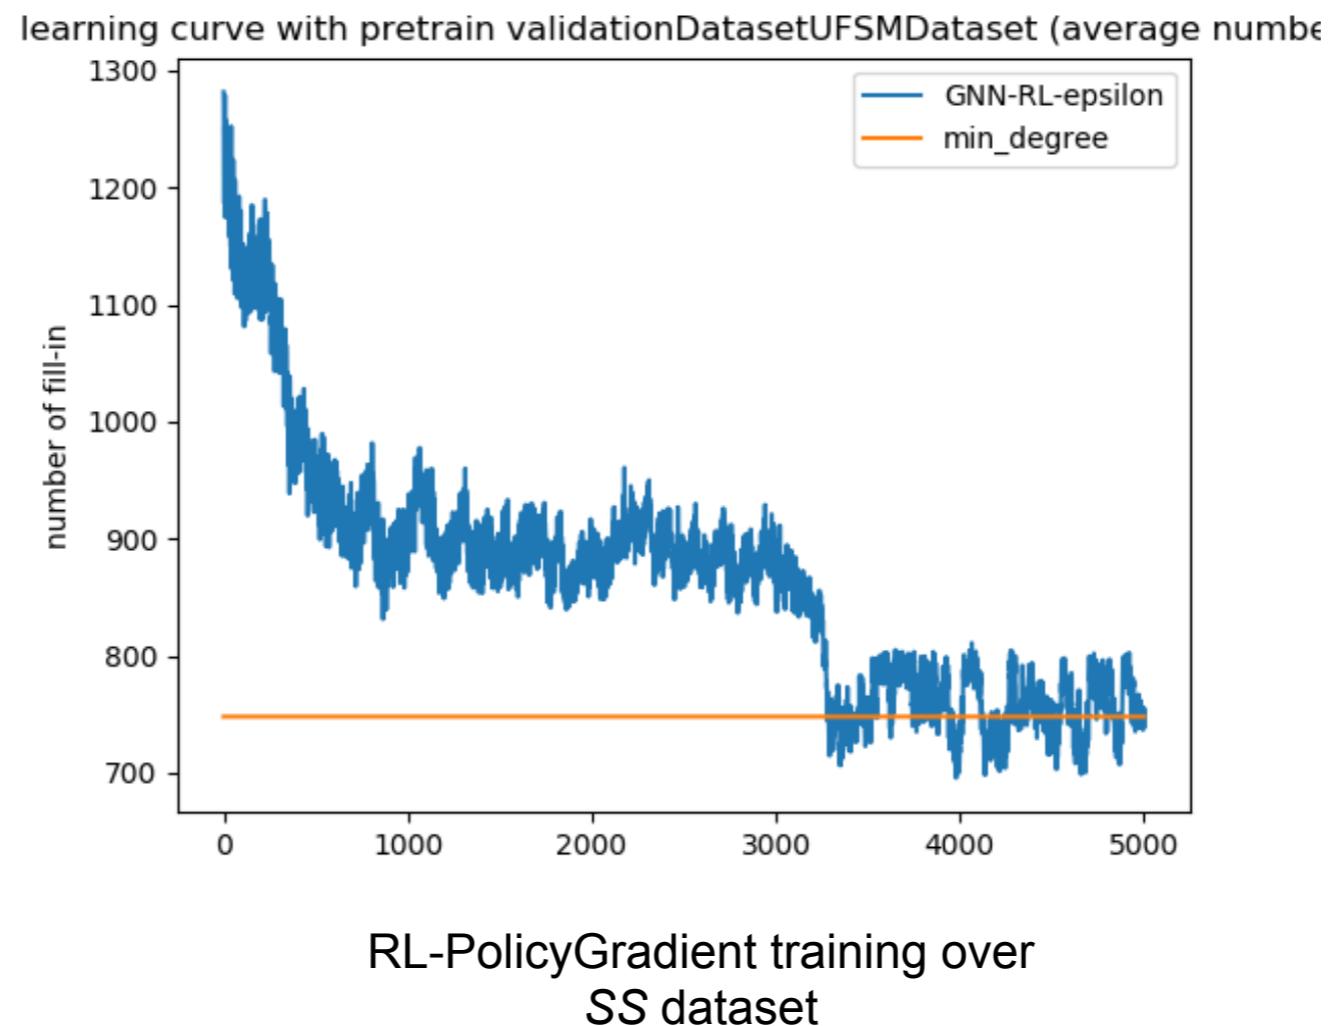


Reinforcement learning framework by Sutton and Barto(1998)

- MDP Formulation:
  - **State space** : set of undirected graphs
  - **Action space** : the set of vertices
  - **Policy** : maps a graph  $G = (V, E)$  to a vertex  $v \in V$  to be eliminated.
  - **Reward/Return( $r$ )**: the number of edges inserted (flexible)

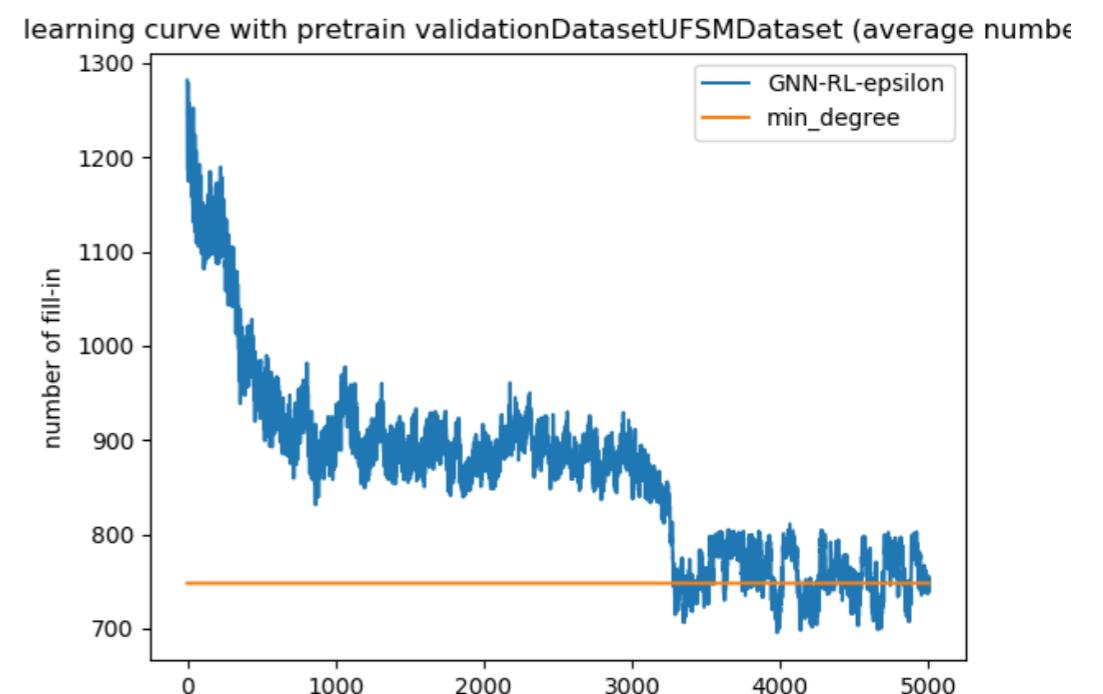
# Preliminary Results:

- RL with Pre-training



# Discussion

- Why does policy model converges to that local min:
  - flat local min landscape of policy function
  - not enough exploration
  - not enough model capacity
- Next step:
  - RL refinement: more exploration to find better local min
  - analysis of model capacity to get theoretical bounds
  - learning a critic function to reduce variance
    - ▶ need to design suitable policy evaluation networks



RL-PolicyGradient training over SS dataset

# Summary

## Applications of GNNs on Combinatorial Optimization:

- Graph/network optimization problems
  - Khalil, et al. "*Learning combinatorial optimization algorithms over graphs.*" NIPS. 2017
  - Kool, et all. "*Attention, learn to solve routing problems!.*" ICLR. 2019.
  - Liu, et all. "*Learning chordal extensions.*" arXiv preprint arXiv:1910.07600 (2019).
- General mixed integer linear programming problems
  - Gasse, et al. "*Exact combinatorial optimization with graph convolutional neural networks.*" Neurips 2019
  - Ding, et al. "*Accelerating Primal Solution Findings for Mixed Integer Programs Based on Solution Prediction.*" arXiv preprint arXiv:1906.09575(2019)

**Thank you for your attention!**